

Next-Generation Foundation Model Architectures

Advanced LLM Research and Next-Generation NLP

Naeemullah Khan

naeemullah.khan@kaust.edu.sa



جامعة الملك عبد الله
للعلوم والتقنية

King Abdullah University of
Science and Technology

KAUST Academy
King Abdullah University of Science and Technology

October 9, 2026

1. Motivation
2. Learning Outcomes
3. Beyond Attention: Recurrent Mixers
 1. Gated Linear Recurrences
 2. One Family of Matrix-State Mixers
 3. Other Families: RWKV, xLSTM, Hyena
4. Hybrid Architectures
5. Mixture-of-Experts at Frontier Scale
6. Efficient Attention
 1. Faster Kernels: FlashAttention-3 and 4
 2. Sliding-Window Attention
 3. Learned Sparse Attention

- 4. Context Parallelism
- 7. Long-Context Architectures
 - 1. Positions and Attention Sinks
 - 2. Training and Evaluating at 1M Tokens
- 8. Scaling Laws and Their Limits
- 9. Emerging Multimodal Architectures
- 10. Summary


The Transformer Was Not Replaced. It Was Re-plumbed.

- ▶ The block from 2017 is still recognisable in every frontier model of 2026.
- ▶ What changed is **where the cost goes**. Three pressures drive almost every design in this lecture:
 - **Memory per token.** Attention keeps every past token in a cache that grows with context.
 - **Compute per token.** Most parameters sit in feed-forward layers that every token pays for.
 - **Context length.** One million tokens is now a product feature, not a research demo.
- ▶ The answers are **hybrids**, **sparsity** and **compression**, not a new architecture.

Frontier MoE models went from about 28% of parameters active per token (Mixtral, 2024) to about 3 to 5% in 2026.

Raschka, "The Big LLM Architecture Comparison" (2025, updated 2026). Jiang et al., "Mixtral of Experts," [arXiv:2401.04088](https://arxiv.org/abs/2401.04088). Kimi K3 and DeepSeek-V4 model configurations (2026).

Today's Map: One Block, Six Questions



By the end of this session, you should be able to:

- ▶ **Relate** decay gates, the delta rule and SSMs, and say why hybrids keep attention.
- ▶ **Compare** Mamba-3, RWKV, xLSTM and Hyena on recall and throughput.
- ▶ **Analyse** MoE routing, load balancing and serving cost at frontier scale.
- ▶ **Explain** how FlashAttention-3 and 4, windows and compressed sparse attention save cost.
- ▶ **Explain** how RoPE scaling, engineered sinks and staged training extend context.
- ▶ **Apply** scaling laws to a training budget, and state where they stop holding.
- ▶ **Describe** the current design choices in natively multimodal models.

Prerequisite knowledge: linear attention, the delta rule and Gated DeltaNet; S4, Mamba and Mamba-2; MoE routing, load balancing and expert parallelism; FlashAttention; block-sparse attention (NSA, DSA); attention sinks; segment recurrence and memory-augmented Transformers; RoPE; effective-context benchmarks; Ring Attention; Kaplan and Chinchilla scaling laws; test-time compute; data curation; multimodal fusion patterns.

Next-Generation Architectures: **Beyond Attention**

Does the sequence mixer have to be attention at all?

- Softmax attention is paid for twice: once in **training compute**, once in **inference memory**.

	Training	State at inference
Classic RNN (LSTM)	sequential	fixed, $O(1)$
Softmax attention	parallel, $O(L^2)$	KV cache, $O(L)$
What we want	parallel, $O(L)$	fixed, $O(1)$

- Albert Gu's picture: a Transformer is a *database* that stores every token. A recurrent model is a *brain* that keeps a compressed state.
- The database gives exact recall. The brain gives constant cost. **This section is about getting both.**

Table after Nayebi and Gormley, CMU 10-423 Lecture 21 (2026). Gu, "On the Tradeoffs of State Space Models and Transformers" (2025), <https://goombalab.github.io/blog/2025/tradeoffs/>.

- ▶ Linear attention keeps a matrix state $S_t \in \mathbb{R}^{d \times d}$ per head, written with key k_t and value v_t , read by query q_t into the output o_t . Gated variants decay the old state by G_t first:

$$S_t = G_t S_{t-1} + v_t k_t^\top, \quad o_t = S_t q_t$$

Model	Decay G_t	Depends on the input?
Linear attention	none, $G_t = I$: the memory never forgets	n/a
RetNet (2023)	one fixed scalar γ per head	no
Mamba-2 (2024)	scalar γ_t	yes
Gated Linear Attention (2023)	per-channel $\text{Diag}(\alpha_t)$	yes

- ▶ Data-dependent gates close most of the language-modelling gap. At 1.3B parameters GLA reaches perplexity 17.22 against 16.85 for a tuned Transformer.
- ▶ Fixed decay did not survive. In a 2025 study, RetNet gives near-zero recall inside hybrids.
- ▶ **What a gate cannot do:** fix one wrong association. The delta rule's erase term $(I - \beta_t k_t k_t^\top)$ can (β_t : write strength).

Yang et al., "Gated Linear Attention Transformers," [arXiv:2312.06635](https://arxiv.org/abs/2312.06635), Table 2. "A Systematic Analysis of Hybrid Linear Attention," [arXiv:2507.06457](https://arxiv.org/abs/2507.06457).

- ▶ A recurrence that is cheap on paper can still be slow on a GPU. Step-by-step updates use no matrix multiplies, so the tensor cores sit idle.

Form	Sequential steps	Cost	Used for
Recurrent	L	$O(Ld^2)$	decoding
Fully parallel	1	$O(L^2d)$	short sequences
Chunkwise, chunk C	L/C	$O(LCd + Ld^2)$	training

- ▶ Inside a chunk: ordinary masked attention. Across chunks: pass the state. In practice C is 64 to 256.
- ▶ More FLOPs than the recurrence, yet faster in wall-clock time. **Hardware use beats FLOP count.**
- ▶ The delta rule needs one more trick: its chain of erase matrices is kept in a compact form (the WY representation).

Yang et al., [arXiv:2312.06635](https://arxiv.org/abs/2312.06635) and [arXiv:2406.06484](https://arxiv.org/abs/2406.06484). Library:
<https://github.com/fla-org/flash-linear-attention>.

Linear attention, gated recurrences, state-space models and the delta rule are all rules for updating a matrix state, read out as $o_t = S_t q_t$.


Model	State update	What it adds
Linear attention	$S_t = S_{t-1} + v_t k_t^\top$	nothing: write only
RetNet	$S_t = \gamma S_{t-1} + v_t k_t^\top$	fixed decay
Mamba-2	$S_t = \gamma_t S_{t-1} + v_t k_t^\top$	input-dependent decay
GLA	$S_t = S_{t-1} \text{Diag}(\alpha_t) + v_t k_t^\top$	per-channel decay
DeltaNet	$S_t = S_{t-1} (I - \beta_t k_t k_t^\top) + \beta_t v_t k_t^\top$	targeted overwrite
Gated DeltaNet	$S_t = S_{t-1} \alpha_t (I - \beta_t k_t k_t^\top) + \beta_t v_t k_t^\top$	decay and overwrite

- Each transition stays **diagonal plus low-rank**. That structure is what keeps chunkwise training possible.

After Yang et al., [arXiv:2406.06484](https://arxiv.org/abs/2406.06484), Table 2, and Yang, Kautz, Hatamizadeh, [arXiv:2412.06464](https://arxiv.org/abs/2412.06464).

Mamba-3 (2026): Designed for Inference First

- **Better discretisation.** A trapezoidal rule replaces the separate conv layer.
- **Complex-valued state.** The state can *rotate*, which scalar decay cannot. It acts as a data-dependent RoPE.
- **MIMO.** More compute per decode step, same state size.



At 1.5B	Mamba-3	Mamba-2
Perplexity	10.24	10.47
Parity task	100%	0.9%

Results stop at 1.5B parameters. No shipped production model used Mamba-3 as of October 2026.

Lahoti et al., "Mamba-3: Improved Sequence Modeling using State Space Principles," ICLR 2026, [arXiv:2603.15569](https://arxiv.org/abs/2603.15569). Figure: Goomba Lab blog, "Mamba-3 Part 1" (2026).

- **Idea (2023):** an RNN that trains in parallel like a Transformer and decodes with constant memory. Open source from the start.
- Its versions climb the same ladder as the update-rule table: fixed decay, then input-dependent decay, then the delta rule.

Version	State	Update rule
RWKV-4 (2023)	vector, $O(d)$	fixed per-channel decay
RWKV-5 and 6 (2024)	matrix	data-dependent decay (v6)
RWKV-7 “Goose” (2025)	matrix	generalised delta rule , vector learning rate

- RWKV-7 separates the key it *removes* from the key it *writes*. The authors prove it recognises all regular languages with a constant number of layers.
- **Status:** the largest RWKV-7 release is 2.9B. No frontier hybrid uses an RWKV mixer.

Peng et al., [arXiv:2305.13048](#) (RWKV-4), [arXiv:2404.05892](#) (Eagle and Finch), [arXiv:2503.14456](#) (RWKV-7). RWKV wiki, architecture history (Nov 2025).

- ▶ Hochreiter's group asks: how far do LSTMs go with modern scale and two repairs? They are **exponential gating** and a **matrix memory** (the mLSTM cell):

$$C_t = f_t C_{t-1} + i_t v_t k_t^\top$$

- ▶ Look familiar? It is gated linear attention plus an **input gate**. There is no erase term.
- ▶ **xLSTM 7B** (2025): trained on 2.3T tokens. About 50% faster generation than Mamba-based models, with quality the authors call comparable to similar-sized LLMs.
- ▶ A scaling study with 672 training runs (ICLR 2026) claims xLSTM is Pareto-better than Llama-style Transformers.

No xLSTM model above 7B exists, and adoption outside the originating lab is limited.

Beck et al., "xLSTM," NeurIPS 2024, [arXiv:2405.04517](#). "xLSTM 7B," [arXiv:2503.13427](#). "xLSTM Scaling Laws," ICLR 2026, [arXiv:2510.02228](#).

- ▶ **Hyena** (2023) mixes tokens with **long convolutions** whose filters come from a small network, interleaved with data-controlled gating. Cost is $O(L \log L)$.
- ▶ On text it matched small GPT models, but recall is its weak point. A 70M attention model beats a 1.4B Hyena on associative recall.
- ▶ It found its home where sequences are **long and untokenised**:

Evo 2 (Arc Institute, 2025): a 40B-parameter genome model trained on 9.3T DNA base pairs, with a context of 1M base pairs at single-nucleotide resolution. Of its 50 layers, 42 are convolutional and 8 are attention.

- ▶ Gu's argument: SSM-like mixers suit raw, high-resolution data. Attention suits tokens that already carry meaning.
- ▶ Even here, **a few attention layers stay**.

Poli et al., "Hyena Hierarchy," [arXiv:2302.10866](https://arxiv.org/abs/2302.10866). Arora et al., [arXiv:2312.04927](https://arxiv.org/abs/2312.04927). Arc Institute et al., "Genome modeling and design across all domains of life with Evo 2," bioRxiv (2025). Evo 2 40B configuration file, <https://github.com/ArcInstitute/evo2>.

Next-Generation Architectures: **Hybrid Architectures**

If a fixed state cannot recall and attention cannot stay cheap, why not use both?

- Keep the cheap recurrent mixer in most layers. Put full attention in a few.

1.3B, 100B tokens	Perplexity ↓	LM eval ↑	Retrieval ↑
Transformer++	16.9	50.9	41.8
DeltaNet	16.9	51.6	34.7
DeltaNet + attention in 2 layers	16.6	51.8	47.9

- A controlled study of 72 models confirms the shape. Language-modelling quality is **flat** in the mixing ratio. Recall **rises steeply** with more attention, and saturates near 3 recurrent layers per attention layer.
- The KV cache shrinks by the same factor, since only attention layers keep one.
- Surprise from the same study: a mixer's standalone quality does not predict its quality inside a hybrid.

Yang, "Linear Transformers for Efficient Sequence Modeling" (talk, 2024), slides 122 to 126. "A Systematic Analysis of Hybrid Linear Attention," [arXiv:2507.06457](https://arxiv.org/abs/2507.06457).

Model	Layer pattern	Total / active	Context
Qwen3.5 (Feb 2026)	3 Gated DeltaNet : 1 attention	397B / 17B	262K
Kimi K3 (2026)	69 KDA + 24 attention layers	2.8T / 104B	1M
Nemotron 3 Super (Mar 2026)	40 Mamba-2 + 8 attention	120.6B / 12.7B	1M
Jamba2 Mini (Jan 2026)	7 Mamba : 1 attention	52B / 12B	262K
Granite 4.0-H (Oct 2025)	9 Mamba-2 : 1 Transformer	32B / 9B	512K
Falcon-H1 (May 2025)	attention Mamba-2 in every block	up to 34B	256K

- ▶ Two camps: **delta-rule** mixers at about 3:1, and **Mamba-2** mixers with 8 to 20% attention layers.
- ▶ Read the ratio, not the label. Most of these models are also sparse MoEs.

Official model cards and configuration files (Qwen, Moonshot AI, NVIDIA, AI21, IBM, TII), 2025 to 2026. Raschka, "LLM Architecture Gallery: Hybrid Attention" (2026).

- Most hybrid claims compare models trained on different data. **Kimi Linear** ran the cleanest test: same data, same recipe, against a full-attention baseline.

	Kimi Linear	Full attention
MMLU-Pro (4K)	51.0	47.2
RULER (128K)	84.3	81.3

- The hybrid is ahead on short-context knowledge and on long-context retrieval at once. In this comparison the recurrent layers cost no quality.
- Serving had to change too. vLLM added a **hybrid cache manager** so a recurrent state and a KV block share the same memory pool.

Kimi Team, "Kimi Linear: An Expressive, Efficient Attention Architecture," [arXiv:2510.26692](https://arxiv.org/abs/2510.26692). vLLM blog, "vLLM Now Supports Qwen3-Next" (Sep 2025).

MiniMax

- ▶ **M1** (2025): 7 linear layers per softmax layer.
- ▶ **M2** (Oct 2025): **full attention in all 62 layers**.
- ▶ Their stated reasons:
 - deficits in multi-hop reasoning at scale that standard benchmarks missed
 - higher sensitivity to numerical precision
 - no mature prefix caching or speculative decoding for linear layers

IBM

- ▶ **Granite 4.0-H** (Oct 2025): Mamba-2 hybrid.
- ▶ **Granite 4.1 and 4.2** (2026): back to a dense Transformer. The 8B dense model matches or beats the larger 4.0 hybrid.

Hybrids are a **contested** default, not a settled one.

MiniMax, "Why Did MiniMax M2 End Up as a Full Attention Model?" (2025). IBM, "Granite 4.1 LLMs: How They're Built" (Apr 2026).

- ▶ **Live designs:** delta-rule linear attention (Qwen, Kimi) and Mamba-2 (NVIDIA, AI21, TII). RWKV, xLSTM and Hyena keep niches.
- ▶ **Open problem 1, recall.** A fixed state has a capacity bound. Hybrids work around it. They do not solve it.
- ▶ **Open problem 2, state tracking.** Scalar decay cannot do parity or permutations. Mamba-3's rotations and RWKV-7's delta rule both attack this.
- ▶ **Open problem 3, systems.** Caching, quantisation and evaluation were all built for attention.

Next. The mixer is only half of the block. The other half, the feed-forward layer, holds most of the parameters.

Next-Generation Architectures: **Mixture-of-Experts**

Must every token use every weight?

Model	Total / active	Active share	Experts	Top- <i>k</i>
Mixtral 8x7B (Dec 2023)	47B / 13B	28%	8	2
DeepSeek-V3 (Dec 2024)	671B / 37B	5.5%	256 + 1 shared	8
Qwen3-235B (Apr 2025)	235B / 22B	9.4%	128	8
Kimi K2 (Jul 2025)	1.04T / 32.6B	3.1%	384 + 1 shared	8
gpt-oss-120b (Aug 2025)	116.8B / 5.1B	4.4%	128	4
DeepSeek-V4-Pro (Apr 2026)	1.6T / 49B	3.1%	384 + 1 shared	6
Kimi K3 (2026)	2.8T / 104B	3.7%	896 + 2 shared	16

- **Sparser:** the active share fell from 28% to between 3 and 5% in about two and a half years.
- **Finer:** from 8 experts to hundreds. Choosing 2 of 8 gives 28 combinations. Choosing 4 of 16 half-size experts gives 1,820, at the same active size.

Official configuration files and technical reports of each model. Combination count: Welleck, CMU 11-711 Lecture 21 (2026).
Kimi Team, "Kimi K2: Open Agentic Intelligence," [arXiv:2507.20534](https://arxiv.org/abs/2507.20534).

Strategy	How it works	Catch
Token choice (top- k)	each token picks its k best experts	load can be very uneven
Expert choice	each expert picks its best tokens	perfect balance, but leaks future tokens
Hash routing	fixed map from token id to expert	no learning at all
Soft MoE	experts process weighted mixes of tokens	not causal

- ▶ For a decoder the choice is forced. If an expert chooses among the tokens of a sequence, each decision depends on *later* tokens.
- ▶ So almost every LLM uses **token choice**, and must then solve load balancing separately.
- ▶ Old ideas return: DeepSeek-V4 uses *hash routing* in its first three layers.

Hashimoto, Stanford CS336 Lecture 4 (2026). Zhou et al., [arXiv:2202.09368](#). Wang et al., [arXiv:2408.15664](#). Puigcerver et al., [arXiv:2308.00951](#). DeepSeek-V4 report, [arXiv:2606.19348](#).

Softmax gate (Mixtral, Qwen3, gpt-oss)

$$g = \text{Softmax}(\text{TopK}(xW_g))$$

- ▶ Experts *compete*. Raising one score lowers the others.

- ▶ Sigmoid scoring pairs naturally with the bias-based balancing on the next slide. DeepSeek-V3, Kimi K2 and K3, and GLM-5 all use the pair.

- ▶ **2026:** DeepSeek-V4 changes the score again, to $\sqrt{\text{softplus}(\cdot)}$. The report gives no ablation for it.

Sigmoid gate (DeepSeek-V3, Kimi, GLM)

$$s_i = \sigma(x^\top e_i), \quad g_i = \frac{s_i}{\sum_{j \in \text{TopK}} s_j}$$


- ▶ Each expert is scored *on its own*, then the selected scores are normalised.

Jiang et al., [arXiv:2401.04088](#). DeepSeek-AI, “DeepSeek-V3 Technical Report,” [arXiv:2412.19437](#). DeepSeek-V4, [arXiv:2606.19348](#). Configuration files of Qwen3, gpt-oss, Kimi K2, Kimi K3 and GLM-5.

- **The dilemma.** A *balance loss* is an extra training term that penalises uneven use of the experts. It fights the language-modelling loss. Too weak: experts collapse. Too strong: quality drops.
- **The fix.** Add a per-expert bias b_i that is used *only to select*:

$$i \in \text{TopK}(\{s_j + b_j\}), \quad g_i \text{ still from } s_i$$

- After each step, nudge b_i up for underloaded experts and down for overloaded ones.
- No gradient ever flows from the balancing rule into the model.



Wang et al., “Auxiliary-Loss-Free Load Balancing Strategy for Mixture-of-Experts,” [arXiv:2408.15664](https://arxiv.org/abs/2408.15664), Figure 2. MaxVio measures how far the busiest expert is above the mean load.


- Why did the bias method help? DeepSeek's own control experiment gives a surprising answer.

Validation loss	1B model	3B model
Auxiliary loss, balanced per <i>sequence</i>	2.258	2.085
Auxiliary loss, balanced per <i>batch</i>	2.253	2.080
Bias method (no auxiliary loss)	2.253	2.080

- The **scope** matters more than the method. Forcing every sequence to use all experts evenly stops experts from specialising. A code file *should* favour code experts.
- Qwen reached the same conclusion from the other side. Computing the balance loss over a **global batch** lifted a 43B model from 54.98 to 57.30 on MMLU.
- Note: frontier models still keep a very small sequence-level loss as a safety net.

DeepSeek-AI, [arXiv:2412.19437](#), Section 4.5.3. Qiu et al., "Demons in the Detail: On Implementing Load Balancing Loss for Training Specialized Mixture-of-Expert Models," [arXiv:2501.11873](#), Table 1.

Do Experts Specialise? An Updated Answer



- **2025 answer:** it depends on training. Mixtral's experts showed no domain specialisation, but with batch-level balancing domain experts appear (above, right). OLMoE, trained from scratch, shows the same.

Figure: Qiu et al., [arXiv:2501.11873](#). Jiang et al., [arXiv:2401.04088](#), Section 5. Muennighoff et al., "OLMoE: Open Mixture-of-Experts Language Models," [arXiv:2409.02060](#).

The Router Is Fragile: Stability and Token Dropping


- ▶ Router logits pass through an exponential. In bfloat16, large logits mean large round-off, and training diverges.
- ▶ **Router z-loss** keeps the logits small:

$$L_z = \frac{1}{B} \sum_{i=1}^B \left(\log \sum_{j=1}^N e^{x_j^{(i)}} \right)^2$$

- ▶ The habit survives in 2026. Experts are stored in **4-bit** formats, while routers stay in higher precision.
- ▶ **Token dropping.** With a fixed expert capacity, another user's request could cost you a token.
- ▶ Modern training is **dropless**. Block-sparse kernels let each expert take any number of tokens, and DeepSeek-V3 drops none in training or inference.

Zoph et al., "ST-MoE: Designing Stable and Transferable Sparse Expert Models," [arXiv:2202.08906](#). Gale et al., "MegaBlocks," [arXiv:2211.15841](#). DeepSeek-AI, [arXiv:2412.19437](#). Hashimoto, Stanford CS336 Lecture 4 (2026).

- ▶ **Question:** if routing works for the FFN, why not for attention heads?
- ▶ **Design:** every layer routes twice. Top-2 of 8 *attention experts*, then top-2 of 8 *MLP experts*. Keys and values are shared. Each attention expert owns its query and output projections.
- ▶ **Result:** 8B total, about 2B active. Trained on 1.25T tokens for **under \$0.1M**, and it beat Llama-2-7B on the Open LLM average (53.0 against 51.0).
- ▶ **Why it matters:** 30,000 H100 GPU-hours is a budget a university lab can reach.
- ▶ **What did not transfer:** none of the frontier configurations in this lecture routes its attention heads. Attention sparsity took a different road, which we meet in the next section.



Shen et al., "JetMoE: Reaching Llama2 Performance with 0.1M Dollars," [arXiv:2404.07413](https://arxiv.org/abs/2404.07413), Figure 1 and Table 3.

- **How fine should experts be?** Define granularity $G = d_{\text{ff}}/d_{\text{expert}}$. The fitted law

$$\mathcal{L}(N, D, G) = c + \left(\frac{g}{G^\gamma} + a \right) \frac{1}{N^\alpha} + \frac{b}{D^\beta}$$

says the compute-optimal G **grows with scale**, and $G = 1$ is almost never optimal.

- **How sparse?** At fixed compute, sparser is better, and the optimal sparsity approaches 1 as models grow.
- **The warning.** At matched loss, sparser models do **worse on reading comprehension**. Loss does not tell the whole story.

An old rule is dead: ST-MoE (2022) found the benefit “quickly diminishes” beyond 256 experts. Kimi K3 uses 896.

Krajewski et al., “Scaling Laws for Fine-Grained Mixture of Experts,” [arXiv:2402.07871](#). Abnar et al., [arXiv:2501.12370](#).
“Hyperparameter Scaling Laws Across MoE Sparsity,” [arXiv:2609.08690](#).

- ▶ **MoE is a throughput play at large batch.** At batch size 1 you pay memory for every expert and save only FLOPs.
- ▶ **Expert parallelism at scale.** DeepSeek serves V3 with experts spread over 144 GPUs in decode, so each MoE layer is an all-to-all across all of them.
- ▶ **The batch must be large.** With 8 of 256 experts active, each expert sees few tokens unless many requests are pooled. In SGLang's tests, overlapping two batches in decode pays off only above 64 to 128 tokens per device.
- ▶ **Low-bit experts.** Experts are most of the weights, and each is used rarely. gpt-oss stores them at 4.25 bits per parameter, so the 120B model fits on one 80 GB GPU.
- ▶ **Co-design.** *LatentMoE* (NVIDIA, Kimi K3) projects each token to a smaller latent space before dispatch, and the experts compute there. The smaller payload pays for more experts and a larger top- k .

DeepSeek, "DeepSeek-V3/R1 Inference System Overview" (2025). SGLang team, large-scale expert parallelism blog (May 2025). OpenAI, gpt-oss model card, [arXiv:2508.10925](https://arxiv.org/abs/2508.10925). "LatentMoE," [arXiv:2601.18089](https://arxiv.org/abs/2601.18089).

- ▶ The newest frontier MoEs activate **3 to 5%** of their weights per token, across hundreds of small experts.
- ▶ Routing is **token choice**, because a decoder must stay causal.
- ▶ Balancing moved from a loss term to a **selection bias**, and from the sequence to the **batch**. That shift is what lets experts specialise.
- ▶ The router is the delicate part: keep its logits small and its precision high.
- ▶ Open design questions: shared experts, dense early layers, and the right score function.

Next. We made the feed-forward block sparse. Can we do the same to attention itself, without giving up exact recall?

Next-Generation Architectures: **Efficient Attention**


Can we keep attention, and its exact recall, but pay less for it?

Idea	What it changes	Still exact attention?
Better kernels	how the same maths runs on the GPU	yes
Local windows	each token sees only its neighbours	no, long range is cut
Learned sparsity	each token sees a <i>chosen</i> subset	no, but the model picks
Context parallelism	one sequence is split across GPUs	yes

- The first idea dates from 2022: **FlashAttention** removed the bottleneck of its day, which was *memory traffic*. The hardware then moved on. So did the bottleneck.

FlashAttention-3: When Softmax Became the Slow Part

- ▶ On an H100, matrix multiply runs at **989 TFLOPS**. Special functions such as exp run at **3.9**. FlashAttention-2 reaches only about 35% of the GPU's peak.
- ▶ **Idea 1, overlap.** Two groups of threads take turns: while one does the softmax, the other runs matrix multiplies. The slow step hides under the fast one.



- ▶ **Idea 2, FP8.** Halve the precision, double the throughput. Per-block scaling and a random rotation of Q and K keep the error $2.6\times$ below naive FP8.
- ▶ **Result:** 1.5 to $2\times$ faster than FlashAttention-2, up to 740 TFLOPS (75% of peak).

Shah et al., "FlashAttention-3: Fast and Accurate Attention with Asynchrony and Low-precision," [arXiv:2407.08608](https://arxiv.org/abs/2407.08608), Figure 1 and Table 3. Dao, FlashAttention-3 blog (2024).

- ▶ From H100 to B200, tensor-core throughput grew about $2.25\times$. The units that compute exp, and shared-memory bandwidth, **did not grow**.
- ▶ So the exponential is now a co-bottleneck in the forward pass. FlashAttention-4 answers with algorithm changes, not only scheduling:
 - compute part of the exp calls in *software*, with a cubic polynomial on idle arithmetic units
 - rescale the running softmax only when the maximum jumps
 - keep intermediate tiles in the new on-chip tensor memory
- ▶ **Result:** 1613 TFLOPS on a B200 (71% of peak), up to $1.3\times$ over cuDNN.

Each GPU generation moves the binding resource: memory traffic, then thread scheduling, then exp throughput. A kernel is never finished.

Zadouri et al., "FlashAttention-4: Algorithm and Kernel Pipelining Co-Design for Asymmetric Hardware Scaling," [arXiv:2603.05451](https://arxiv.org/abs/2603.05451). Dao, FlashAttention-4 blog (2026).

- ▶ Kernels do not shrink the KV cache. The simplest way to do that: **most layers look only at recent tokens.**

Model	Pattern (local : global layers)	Window
Mistral 7B (2023)	all local	4096
Gemma 2 (2024)	1 : 1	4096
Gemma 3 (2025)	5 : 1	1024
gpt-oss (2025)	1 : 1, with learned attention sinks	128
Llama 4 (2025)	3 : 1, local layers attend within fixed chunks	8192

- ▶ The trend: **windows shrink, and a few global layers carry the long range.**
- ▶ Gemma 3: KV memory overhead at 32K tokens falls from about 60% to under 15%.
- ▶ The same shape as the hybrids. Cheap layers do most of the work, and a few expensive ones keep the capability.



Jiang et al., [arXiv:2310.06825](#). Gemma Team, [arXiv:2408.00118](#) and [arXiv:2503.19786](#), Figure 5. OpenAI, [arXiv:2508.10925](#). Hugging Face, Llama 4 release blog (2025).

- ▶ **A common misreading.** Stacking k layers of window W gives a theoretical span of $W \times k$. That is a receptive field, not reliable long-range recall.
- ▶ **Evidence.** In MiniMax's ablation, full attention scores 90 on RULER at 128K. The sliding-window variant scores 72.
- ▶ **Memory is not latency.** Mistral dropped sliding windows in a later model: they reduce memory but do not necessarily reduce inference latency.
- ▶ Fixed patterns were tried long ago (Sparse Transformer, Longformer, BigBird, 2019 to 2020). Only the local window survived.

A fixed pattern decides where to look before reading the input. What if the model *chose* which distant tokens matter?

Raschka, "MiniMax M2 Technical Report Notes" (May 2026) and "The Big LLM Architecture Comparison" (2025, updated 2026). Child et al., [arXiv:1904.10509](#). Beltagy et al., [arXiv:2004.05150](#). Zaheer et al., [arXiv:2007.14062](#).

- ▶ In DeepSeek Sparse Attention (V3.2) a lightning indexer picks the top 2,048 tokens, but the indexer itself still scans every token. V4 shrinks the sequence *before* scoring.
- ▶ Two layer types alternate through the network:
 - **CSA**: compress keys and values $4\times$, then keep the top- k compressed entries. A sparse, fine view.
 - **HCA**: compress $128\times$ and attend to all of it. A dense, coarse view.
- ▶ Every layer adds a window of 128 recent uncompressed tokens.
- ▶ At 1M tokens, a CSA layer in V4-Pro reads about 1,150 entries. An HCA layer reads about 8,300.
- ▶ The cache shrinks with it: DeepSeek-V4.1-Flash reports about 890 bytes of KV per token.



DeepSeek-AI, “DeepSeek-V4,” [arXiv:2606.19348](https://arxiv.org/abs/2606.19348), Figure 1 (right) and Section 2.3. Entry counts derived from the published hyperparameters ($k = 1024$, window 128, $2^{20}/128$). DeepSeek-AI, [arXiv:2609.19969](https://arxiv.org/abs/2609.19969).

- ▶ **June 2026.** Seven months after returning to full attention for M2, MiniMax ships M3 with its own block-sparse attention: $14.2\times$ faster prefill and $7.6\times$ faster decoding at 1M tokens, on par with the dense baseline.
- ▶ **What changed?** Not the idea. The *kernels* (selection matched to GPU blocks) and the *training recipes* (adapt from a dense checkpoint).
- ▶ The frontier in mid-2026 is the indexer itself. GLM-5.2 reuses one indexer result across four layers and removes 75% of the indexer's compute.
- ▶ No training budget? **Training-free** methods (Quest, MInference, DuoAttention) select KV entries at inference time on an unmodified model.

Lai et al., “MiniMax Sparse Attention,” [arXiv:2606.13392](#). Bai et al., “IndexCache,” [arXiv:2603.12201](#). Tang et al., [arXiv:2406.10774](#).

- ▶ **A plain ring is unbalanced under a causal mask.** When Ring Attention splits one sequence across devices, the device holding the last tokens has far more work than the one holding the first. *Striped Attention* interleaves tokens across devices to even the load.
- ▶ **What Meta did for Llama 3.** No ring at all. Each device simply gathers all keys and values. Two reasons: with grouped-query attention the KV tensors are small, and the gather is $O(S)$ under an $O(S^2)$ computation.

Setup	Sequence length	GPU utilisation (MFU)
Llama 3, 16,384 GPUs, 16-way context parallel	131,072	38%
torchTitan, Llama3-8B on 32 H100, all-gather	1M	25.7%
torchTitan, same, ring-style all-to-all	1M	12.1%

- ▶ **Open question:** how to split learned sparse attention, where the tokens a query needs are not known in advance.

Brandon et al., [arXiv:2311.09431](https://arxiv.org/abs/2311.09431). Meta, “The Llama 3 Herd of Models,” [arXiv:2407.21783](https://arxiv.org/abs/2407.21783), Table 4. PyTorch forum, “Training long-context LLMs with 1M sequence length using context parallel” (Jan 2025).

- ▶ **Kernels** (FlashAttention-3 and 4) give exact attention at 70 to 75% of GPU peak. They follow the hardware, generation by generation.
- ▶ **Windows** cut the KV cache cheaply, but a few global layers must remain.
- ▶ **Learned sparse attention** became a production default within 2025 to 2026: a cheap scorer, a top- k selection, and a local window.
- ▶ **Context parallelism** is how million-token sequences are trained at all.

Next. We can now afford a million tokens. Affording them and *using* them are different problems.

Next-Generation Architectures: **Long Context**

We can now afford a million tokens. Can the model actually use them?

1. **Positions.** The model meets position indices it never saw in training. With RoPE, that means rotation angles outside the trained range.
2. **Attention itself.** Trained heads rely on habits, such as dumping spare attention on the first token. Change the window and those habits fail.
3. **Finding things.** Even when nothing breaks, retrieving the right fact among a million tokens is hard. And we need honest ways to measure it.


This section follows that order: **extend the positions**, **engineer the sinks**, then **train and evaluate** at 1M tokens.

Why RoPE Fails Beyond Its Training Length

- ▶ RoPE rotates each pair of dimensions at its own frequency. Dimension d has wavelength

$$\lambda_d = 2\pi b^{2d/|D|}$$

- ▶ **Fast** dimensions complete many rotations inside the training length L . They have seen every angle.
- ▶ **Slow** dimensions never finish one rotation. Past L they produce angles the model has **never seen**.
- ▶ **Position Interpolation:** squeeze new positions into the old range, $m \mapsto mL/L'$. It works after a short fine-tune, but it also squeezes the fast dimensions, which blurs *local* order.



Schematic from the YaRN ramp with $\alpha = 1$, $\beta = 32$,
 $s = 8$.

Chen et al., "Extending Context Window of Large Language Models via Positional Interpolation," [arXiv:2306.15595](#). Peng et al., "YaRN," [arXiv:2309.00071](#).

- ▶ Let $r(d) = L/\lambda_d$ be the rotations a dimension completes in training. YaRN blends per dimension:

$$\theta'_d = (1 - \gamma(r)) \frac{\theta_d}{s} + \gamma(r) \theta_d, \quad \gamma(r) = \begin{cases} 0 & r < \alpha \\ 1 & r > \beta \\ \frac{r - \alpha}{\beta - \alpha} & \text{otherwise} \end{cases}$$

- ▶ **Fast dimensions** are left alone. **Slow dimensions** are fully interpolated. The rest are blended.
- ▶ Longer contexts also flatten attention, so YaRN adds a temperature, $\sqrt{1/t} = 0.1 \ln s + 1$.
- ▶ **Cost:** about 400 fine-tuning steps, against 1,000 for Position Interpolation.

Still in use in 2026. gpt-oss, Qwen3-Next and DeepSeek-V4 all ship with YaRN scaling in their configurations.

Peng et al., "YaRN: Efficient Context Window Extension of Large Language Models," [arXiv:2309.00071](https://arxiv.org/abs/2309.00071). Model cards and configurations of gpt-oss, Qwen3-Next and DeepSeek-V4.

- ▶ **Search it.** *LongRoPE* finds a separate scale for every dimension by evolutionary search, and reaches 2M tokens. *LongRoPE2* adds one insight: slow dimensions are under-trained, so the cut-off sits earlier than theory says. On Llama-3-8B at 128K it scores 82.0 on RULER with **10B** tokens of training. Meta used 800B.
- ▶ **Train it.** Frontier labs now build extension into pre-training. They raise the RoPE base (Llama 3 uses 500,000) and lengthen sequences in stages.
- ▶ **Remove it.** A growing group uses **no positional encoding** in the global-attention layers, and leaves order to local or recurrent layers:
 - Llama 4 interleaves NoPE layers with RoPE layers
 - Nemotron 3 and Kimi Linear: the recurrent layers already encode order
 - DeepSeek-V4 and Qwen3-Next apply RoPE to only a fraction of the dimensions


Ding et al., “LongRoPE,” [arXiv:2402.13753](#). Shang et al., “LongRoPE2,” [arXiv:2502.20082](#), Table 2. Meta, [arXiv:2407.21783](#) and Llama 4 blog (2025). NVIDIA, [arXiv:2512.20856](#).

- Sinks get **stronger** with longer training context and larger models, so recent designs build in a “do nothing” option.
- **Fix A: give it a real null slot.** Add a learned term to the denominator (gpt-oss, DeepSeek-V4):

$$a_{ij} = \frac{e^{z_{ij}}}{\sum_k e^{z_{ik}} + e^{z'}}$$

with z_{ij} the attention logits and z' a learned logit per head.

- **Fix B: let the head switch off.** Gate the attention output, $Y' = Y \odot \sigma(XW_{\theta})$ (Qwen).



With output gating, mean first-token attention falls from 0.467 to 0.048.

Barbero et al., “Why do LLMs attend to the first token?” [arXiv:2504.02732](#). DeepSeek-V4, [arXiv:2606.19348](#), Eq. 27. Figure: Qiu et al., “Gated Attention for Large Language Models,” [arXiv:2505.06708](#).

How Models Are Actually Trained to 1M Tokens

Model	Length schedule	Tokens spent on long stages
Llama 3 (2024)	6 stages, 8K to 128K	about 800B
MiniMax-01 (2025)	128K, then 512K, then 1M	300B + 32B + 26B
Nemotron 3 Super (2026)	8K, then straight to 1M	34B + 17B
DeepSeek-V4 (2026)	4K, 16K, 64K, then 1M	not reported per stage

- ▶ **The last stage is cheap.** A few tens of billions of tokens at 1M, after trillions at short length.
- ▶ **Gate each stage.** Llama 3 moved on only when short-context quality had recovered *and* needle retrieval was perfect at the new length.
- ▶ **Needle tests saturate early.** MiniMax found them solved long before real long-context ability appeared, and tracked harder tasks instead.

Meta, [arXiv:2407.21783](#). MiniMax, [arXiv:2501.08313](#), Table 6. NVIDIA, “Nemotron 3 Super,” [arXiv:2604.12374](#). DeepSeek-AI, [arXiv:2606.19348](#).

The 2026 Picture: What Happens Between 128K and 1M

Model	to 128K	to 1M	Decay
Claude Opus 4.6	77.1	70.6	8.5%
Gemini 3.1 Pro	77.8	68.5	12.0%
GPT-5.5	74.6	67.8	9.2%
DeepSeek-V4-Pro	71.6	59.1	17.4%
Qwen3.5-397B	71.0	54.4	23.3%
GLM-5	64.3	44.3	31.1%

ATLAS score, averaged over lengths from 8K up to the stated limit.

- ▶ An **independent** harness over 26 models: the average score drops **24.3%** when the range extends from 128K to 1M.
- ▶ Retrieval and question answering lose the most at 1M.
- ▶ Closed frontier models decay by 8 to 14%. Open models by 17 to 60%.
- ▶ **Beware vendor tables.** Labs report the same benchmark in different formats, and their numbers for each other's models disagree.

A 1M-token window is real. Uniform quality across that window is not there yet.

Huang et al., "ATLAS," [arXiv:2605.28079](https://arxiv.org/abs/2605.28079), Tables 10 and 11 (May 2026).

- ▶ **Positions:** scale RoPE by frequency band (YaRN), or drop positional encoding from global layers altogether.
- ▶ **Sinks:** softmax needs a “do nothing” option. Modern models provide one on purpose.
- ▶ **Training:** stage the length, gate each stage, and spend little at the longest length.
- ▶ **Evaluation:** trust independent, reasoning-heavy benchmarks over needle tests and vendor tables.

Next. Every design so far changes what a parameter or a token costs. How should a fixed budget be divided between model size, data and inference?


Next-Generation Architectures: **Scaling Laws and Their Limits**

Given a fixed budget, how big should the model be, and how long should it train?

- ▶ **Notation.** Loss is a power law in parameters N and tokens D , with compute $C \approx 6ND$. The compute-optimal model size grows as $N^* \propto C^a$.
- ▶ **Correction 1.** Chinchilla's own parametric fit was wrong. Taken literally it implies about 70 tokens per parameter. A careful refit recovers 20, and shows the reported confidence intervals were implausibly tight.
- ▶ **Correction 2.** Kaplan and Chinchilla never disagreed about Transformers. The gap is a *measurement artefact*, as the next slide shows.

“20 tokens per parameter” is a property of one training setup. It is a useful heuristic, not a law.

Hoffmann et al., [arXiv:2203.15556](#). Besiroglu et al., “Chinchilla Scaling: A replication attempt,” [arXiv:2404.10102](#). Pearce and Song, [arXiv:2406.12907](#).



- ▶ Over 900 runs reproduce Kaplan's setup (0.84 here, 0.88 in Kaplan et al.), then close the gap.
- ▶ **Three corrections:** count the output layer's compute (the biggest), scale warmup with model size, tune the optimiser per scale.
- ▶ Cosine decay barely moves a , and the final run uses a constant learning rate: decay is *not* needed for the 0.5 law.
- ▶ **Lesson for your own fits:** exponents are fragile to protocol. Small models and untuned hyperparameters bias them.

Porian et al., "Resolving Discrepancies in Compute-Optimal Scaling of Language Models," [arXiv:2406.19146](https://arxiv.org/abs/2406.19146), Figure 1 and Table 1 (RefinedWeb).

Model	Tokens / param
GPT-3	2
Chinchilla	20
Llama 2 70B	29
Mistral 7B	110
Llama 3 70B	215
Llama 3 8B	about 1,900

Last row derived from 15T tokens and 8B parameters.

- ▶ Chinchilla minimises **training** compute for a target loss. A deployed model also pays for every token it ever generates.
- ▶ If inference demand is large, the best choice is a **smaller model trained far longer**. Quality keeps improving up to 10,000 tokens per parameter.
- ▶ Meta's own account: the Chinchilla-optimal budget for an 8B model is about 200B tokens. Llama 3 8B kept improving log-linearly up to 15T.
- ▶ Flagships stay near compute-optimal. **Small models are over-trained by 100× or more.**

Table: Hashimoto, Stanford CS336 Lecture 9 (2026). Sardana et al., "Beyond Chinchilla-Optimal," [arXiv:2401.00448](https://arxiv.org/abs/2401.00448). Meta, "Introducing Meta Llama 3" (Apr 2024).

- ▶ Scaling laws assume fresh data. What if you must repeat it?
- ▶ **Repetition has a budget.** Up to about 4 epochs is almost as good as new data. The value of further epochs decays exponentially:

$$D' = U_D + U_D R_D^* (1 - e^{-R_D/R_D^*}), \quad R_D^* \approx 15$$

where U_D is unique tokens and R_D the number of repetitions.

- ▶ **Responses:** stronger regularisation, synthetic data for maths and code, and compute spent on reinforcement learning instead of more text.

Sutskever, NeurIPS 2024: “We’ve achieved peak data and there’ll be no more. There’s only one internet.”

Muennighoff et al., “Scaling Data-Constrained Language Models,” [arXiv:2305.16264](#). Kim et al., [arXiv:2509.14786](#). Quote as reported by The Verge (13 Dec 2024).

- ▶ Kaplan's finding still holds at scale: for a fixed parameter count, loss depends only weakly on shape. So shape is set by convention and by hardware.

Choice	Convention	Recent outliers
FFN width / model width	4, or 8/3 with gated units	Gemma 2 uses 8
Model width / depth	100 to 200	Gemma 3: 87. Gemma 4: 61
Heads $\times$ head dim / width	1	Qwen 3.5: 1.2
Vocabulary	100K to 250K in production	Gemma 4: 262K

- ▶ **Where shape does matter:**
 - *Small models.* Below a billion parameters, deep and thin beats shallow and wide (MobileLLM).
 - *Vocabulary.* The optimal size grows with compute. By one estimate, Llama 2 70B's 32K vocabulary should have been at least 216K.

Table: Hashimoto, Stanford CS336 Lecture 3 (2026). Liu et al., "MobileLLM," [arXiv:2402.14905](#). Tao et al., "Scaling Laws with Vocabulary," [arXiv:2407.13623](#).

- **Precision.** Treat bit-width as a third resource. Low precision acts like a reduction in effective parameters:

$$N_{\text{eff}} = N (1 - e^{-P_w/\gamma_w}) (1 - e^{-P_a/\gamma_a}) (1 - e^{-P_{kv}/\gamma_{kv}})$$

Two consequences. About 7 to 8 bits is compute-optimal for training. And **over-trained models quantise worse**, which hits small deployed models hardest.

- **Architecture.** Hybrids, MoE and multimodal models get their own fitted laws. The exponents stay close to the dense Transformer's.
- **Emergence.** Are sudden jumps in ability real? Two answers that fit together:
 - with continuous metrics, most “emergent” curves become smooth
 - abilities appear when the *pre-training loss* crosses a threshold, whatever the model size
- So the working recipe is two-stage: predict loss from compute, then task score from loss.

Kumar et al., “Scaling Laws for Precision,” [arXiv:2411.04330](#). Schaeffer et al., [arXiv:2304.15004](#). Du et al., [arXiv:2403.15796](#). MiniMax, [arXiv:2501.08313](#).

- ▶ **Compute keeps growing.** Frontier training compute has risen about $5\times$ per year since 2020. The largest estimated run so far is Grok 4, at about 5×10^{26} FLOP.
- ▶ **Pre-training alone pays less.** GPT-4.5 used roughly ten times GPT-4's compute for incremental gains. GPT-5 then used *less* training compute than GPT-4.5: scaling post-training on a smaller model gave better returns.

Be careful with “scaling has hit a wall”. No lab has published a frontier-scale loss curve that shows the power law breaking.

Epoch AI, trends dashboard (Feb 2026) and “Why GPT-5 used less training compute than GPT-4.5” (Sep 2025). Lambert, Interconnects (Feb 2025).

- ▶ The compute-optimal rule survives, but **its constants depend on protocol**. Fit your own.
- ▶ Deployed models are deliberately **over-trained**, because inference dominates lifetime cost.
- ▶ **Data** is now a binding constraint: repeating it pays for about 4 epochs, then decays.
- ▶ Compute has three destinations: **pre-training, post-training and inference**. Each has a different curve.

Next. All of this assumed text. The last section asks what changes when the model must also see, hear and draw.

Next-Generation Architectures: **Multimodal Foundation Models**

What changes when the model must also see, hear and draw?

Model	Backbone	How vision is built in
Kimi K2.5 (Jan 2026)	MoE, 1.04T / 32B active	a 400M native-resolution ViT, with 15T tokens of joint vision and text pre-training
Qwen3.5 (Feb 2026)	3:1 hybrid MoE, 397B / 17B	one checkpoint replaces the separate Qwen3-VL line
Gemma 4 (Jul 2026)	2.3B to 31B, dense and MoE	a 550M ViT in most sizes

- ▶ Until 2025, most labs shipped a text model and a separate “-VL” variant. Each of these three ships **one natively multimodal checkpoint**.
- ▶ Note the backbones. They are the hybrid and MoE designs from earlier today. The novelty is in the backbone and in token budgets, not in a new fusion trick.

Kimi Team, “Kimi K2.5,” [arXiv:2602.02276](#). Qwen3.5-397B-A17B model card. Gemma Team, “Gemma 4,” [arXiv:2607.02770](#).

Terms. *Early fusion*: image patches enter the first layer, with no vision encoder. *Encoder-based*: a pre-trained vision encoder feeds the language model (Shukor et al. call it late fusion).

A vendor ablation (Kimi K2.5)

- ▶ With a fixed total budget of vision and text tokens, adding vision **early in training, at only 10%** of the mix, beats adding it **late at 50%**.
- ▶ The ratio matters little. The timing matters.
- ▶ A bonus: reinforcement learning on *vision* tasks also raised *text* scores, by about 2 points on GPQA-Diamond.

An academic scaling study (Apple)

- ▶ 457 models, trained from scratch. At equal compute, early fusion is **on par** with the encoder-based design, and slightly better at small budgets.
- ▶ Both follow almost the same scaling exponent as text-only models:

	$L \propto C^x$
Early fusion	-0.0492
Encoder-based	-0.0494
GPT-3 (text)	-0.048

An encoder's real advantage is *reuse* of pre-trained parts. Enough native training shrinks it.

Kimi Team, [arXiv:2602.02276](https://arxiv.org/abs/2602.02276), Tables 1 and 2. Shukor et al., "Scaling Laws for Native Multimodal Models," ICCV 2025, [arXiv:2504.07951](https://arxiv.org/abs/2504.07951).

- ▶ If one backbone handles every modality, should some weights be reserved for each?
- ▶ **Mixture-of-Transformers (Meta).** Give each modality its own feed-forward, attention and normalisation weights, but keep *one global self-attention* over the whole sequence. Routing is by modality, with no learned router. It matches a dense Chameleon-style model at **55.8%** of the FLOPs.
- ▶ **Learned MoE (Apple).** Let the router decide. Routing that ignores modality labels beats routing that uses them. Experts still specialise by modality on their own, mostly in the first and last layers.
- ▶ These agree more than they differ:

Specialise the capacity. Keep the attention global. It is the same division of labour as MoE in a text model: experts differ, mixing is shared.

Liang et al., “Mixture-of-Transformers,” [arXiv:2411.04996](#). Shukor et al., [arXiv:2504.07951](#).

- ▶ An image can cost over a thousand tokens. A video costs that per frame. With long context expensive, **visual-token compression** became a design axis of its own.
 - **Budgets:** Gemma 4 lets the caller choose 70 to 1,120 tokens per image.
 - **Temporal pooling:** Kimi K2.5 averages groups of four video frames at the patch level, and fits about $4\times$ longer video.
 - **Optical compression:** DeepSeek-OCR reads a page of text from an image with about 97% precision at under $10\times$ fewer tokens than the text itself. At $20\times$ it falls to about 60%.
- ▶ **Does the encoder have to exist?** Gemma 4's 12B model has none. Raw 48×48 pixel patches go through a single linear projection into the LLM. Audio chunks are projected the same way.

No published ablation yet compares encoder-free and encoder-based models at production scale. The other Gemma 4 sizes keep a ViT.

Gemma Team, [arXiv:2607.02770](https://arxiv.org/abs/2607.02770). Kimi Team, [arXiv:2602.02276](https://arxiv.org/abs/2602.02276). Wei et al., "DeepSeek-OCR: Contexts Optical Compression," [arXiv:2510.18234](https://arxiv.org/abs/2510.18234).

How images are generated	Example	GenEval $\uparrow$	MMMU $\uparrow$
Next-token prediction over image tokens	Emu3.5 (34B)	0.86	n/a
Autoregressive text, flow-matching images	BAGEL (14B)	0.82	55.3
Masked discrete diffusion for everything	LLaDA2.0-Uni (16B)	0.89	50.1
Multimodal LLM with a diffusion head	InternVL-U (4B)	0.85	54.7

- ▶ **No design has won.** Four very different mechanisms land within a few points of each other. Sizes and settings differ, so read the numbers as indicative.
- ▶ **Do the two tasks help each other?** A 2026 analysis says yes at the level of representations, and no if both are forced through identical computation. Share the knowledge. Decouple the pathway.

Scores as reported in: ByteDance Seed, "BAGEL," [arXiv:2505.14683](#). [arXiv:2603.09877](#) (InternVL-U). [arXiv:2604.20796](#) (LLaDA2.0-Uni). [arXiv:2601.02204](#) (NextFlow, for Emu3.5). [arXiv:2609.01607](#).

- ▶ The flagship of 2026 is **one natively multimodal model**, built on the same hybrid and MoE backbones as text models.
- ▶ **Bring vision in early.** Adding it early in training beats adding it late, and an encoder-free model matches an encoder-based one at equal compute.
- ▶ **Specialise capacity, share attention.** Modality-specific experts with global mixing.
- ▶ **Visual tokens are a budget** to be compressed and chosen, like context length.
- ▶ **Unified generation is unsettled.** Autoregression, flow matching and discrete diffusion all remain in play.

Closed labs disclose almost no architectural detail. What we can teach rigorously comes from open-weight reports.

Pressure	The 2026 answer	Still open
KV cache grows with context	hybrids: recurrent mixers plus a few attention layers	recall and state tracking in a fixed state
Every token pays for every weight	MoE with 3 to 5% of weights active	routing score, shared experts, serving at small batch
Attention is quadratic	FlashAttention-3 and 4, windows, compressed sparse attention	the cost of the selector itself
Context beyond training length	YaRN, NoPE layers, explicit sinks, staged training	uniform quality across 1M tokens
Finite compute and data	over-train small models, repeat data up to about 4 epochs, shift compute to post-training	laws for post-training and for verifiers
Inputs beyond text	one natively multimodal model, vision added early in training	unified generation, encoder-free inputs

The common thread: **do the expensive, exact operation rarely, and the cheap, approximate one everywhere else.**

- ▶ Hybrids and sparse attention both claim the long-context frontier. Under matched training and serving budgets, **which wins on multi-hop reasoning**, and why?
- ▶ Can a recurrent state be made to **grow on demand**, so capacity follows the input instead of being fixed in advance?
- ▶ What is the right **benchmark** for a million-token context, now that needle tests are saturated and vendor tables disagree?
- ▶ Sparse models lose on reading comprehension at matched loss. **What does loss fail to measure?**
- ▶ Is there a scaling law that jointly allocates compute between **pre-training, reinforcement learning and inference**?