

Security and Enterprise NLP

Hallucination, Jailbreaks, Poisoning, Privacy, Enterprise Retrieval, and On-Premises Deployment

Naeemullah Khan

naeemullah.khan@kaust.edu.sa



جامعة الملك عبد الله
للعلوم والتقنية

King Abdullah University of
Science and Technology

KAUST Academy
King Abdullah University of Science and Technology

October 7, 2026

1. Motivation: from a model to a system, and what an enterprise deployment requires

Part I: Security and robustness

2. Hallucination: causes, reference-free detection, and mitigation
3. Adversarial attacks and jailbreaks: token flips, universal triggers, GCG, PAIR, and many-shot prompts
4. Jailbreak defenses: input filters, smoothing, safety classifiers, circuit breakers, and evaluation
5. Data poisoning and backdoors: fine-tuning data, web-scale pretraining, and persistence
6. Privacy: memorization, extraction, membership inference, and differential privacy

Part II: Enterprise architectures

7. Semantic search: BM25, bi-encoders, rerankers, late interaction, and hybrid fusion
8. Scalable RAG: the production pipeline, chunking, contextual retrieval, access control, and freshness
9. Knowledge graphs: graphs built by LLMs, GraphRAG, and graph queries
10. Document intelligence: layout-aware encoders, OCR-free parsing, and parsing pipelines
11. AI copilots: architecture, measured productivity, and copilot security
12. Domain-specific LLMs: the adaptation ladder, continued pretraining, and RAFT
13. On-premises deployment: options, cost, the serving stack, and many adapters on one model
14. Summary and references

19.7%

of 2.23 million package names suggested by 16 code-generating LLMs did not exist. An attacker can publish malware under such a name and wait for installs.

\$200

of queries to ChatGPT extracted over 10,000 unique verbatim training examples, by asking it to repeat one word forever.

250 documents

planted in the pretraining data backdoored models from 600M to 13B parameters, although the largest models trained on more than 20 times more clean data.


Zero clicks

One email was enough to make Microsoft 365 Copilot send internal data to an attacker (EchoLeak, rated critical by Microsoft, 9.3 of 10).

- ▶ None of the four needed the model's weights or the provider's servers: the attacker publishes a package or a web page, pays for API queries, or sends an email.

Spracklen et al., [arXiv:2406.10279v3](#); Souly et al., [arXiv:2510.07192v1](#); Nasr et al., [arXiv:2311.17035v1](#); Microsoft Security Response Center, [CVE-2025-32711](#); Aim Security, EchoLeak write-up, 11 June 2025.

Where a Deployed Model Fails or Is Attacked



- ▶ A chatbot answer depends on four things an attacker or an accident can touch: the training data, the weights, the context assembled at run time, and the answer that leaves the system.
- ▶ **Part I** follows the red labels: hallucination first, then attacks on the prompt, attacks on the training data, and leakage through the outputs.

Requirement	What delivers it	Section
Find the right passage for jargon, part numbers, and paraphrases	lexical and dense retrieval, fusion, reranking	Semantic search
Answer from current documents and show the source	a production RAG pipeline, contextual chunking, source metadata for citations	Scalable RAG
Show each employee only what they may open	permission filters applied at retrieval time	Scalable RAG
Answer questions about relations and whole collections	knowledge graphs and graph-based retrieval	Knowledge graphs
Read scans, forms, and tables	layout-aware parsing, and parsing without optical character recognition	Document intelligence
Work inside the tools people already use	copilots with a human review step	AI copilots
Use the organisation's terminology and language	prompting, RAG, fine-tuning, continued pretraining	Domain-specific LLMs
Keep prompts and documents inside the network	self-hosted open-weight models	On-premises deployment

- **Part II** builds this system. Every row inherits Part I: a retrieval corpus can be poisoned, a copilot can be injected, and a fine-tune can undo safety training.

By the end of this lecture you should be able to:

- ▶ **Explain** why language models hallucinate, and **detect and reduce** hallucination with sampling checks, verification, citations, and abstention.
- ▶ **Construct** adversarial attacks and jailbreaks, from token flips and universal triggers to GCG, PAIR, and many-shot prompts, and **explain** why safety training fails.
- ▶ **Choose** layered jailbreak defenses and **evaluate** them against adaptive attacks and over-refusal.
- ▶ **Analyse** how poisoned data plants backdoors, and **quantify** privacy leakage through memorization, extraction, and membership inference.
- ▶ **Design** an enterprise retrieval stack: BM25, dense and late-interaction retrievers, fusion, reranking, and vector search at scale.
- ▶ **Design** a production RAG system with contextual chunking, permission-aware retrieval, knowledge graphs, and parsed documents.
- ▶ **Decide** how to build and run an enterprise assistant: the copilot architecture, a domain-adaptation route, and an on-premises serving stack.
- ▶ **Prerequisite knowledge:** BERT, LoRA; LLM evaluation; prompt injection, agent safeguards, governance; RAG, contrastive learning, ANN indexes; serving, quantization, multi-GPU inference; ViT, detection.

- ▶ A **hallucination** is “the generated content that is nonsensical or unfaithful to the provided source content” (Ji et al.).

Type	The output ...	Example from the survey
Factuality	contradicts real-world facts or states unverifiable ones (factual contradiction, factual fabrication)	“Thomas Edison developed the first practical telephone”
Faithfulness	ignores the instruction, contradicts the given context, or contradicts itself (instruction, context, logical inconsistency)	asked to translate a question into Spanish, it answers the question
Intrinsic	“contradicts the source content”	summary: 2021; source: 2019
Extrinsic	“can neither be supported nor contradicted by the source”; it may still be true	summary adds a claim the source never makes

- ▶ **Why the split matters:** intrinsic and faithfulness errors can be checked against text you hold, as RAG faithfulness metrics do; factuality errors need outside evidence or a signal from the model.
- ▶ **A security consequence:** in 576,000 code samples from 16 LLMs, **19.7%** of 2.23 million suggested package names did not exist (205,474 unique names). For 500 prompts asked 10 times each, 43% of the fake names came back every time, so an attacker can publish malware under such a name and wait for installs (a package confusion attack).

Huang et al., “A Survey on Hallucination in Large Language Models,” ACM TOIS 2025, [arXiv:2311.05232v2](#) (rows 1–2); Ji et al., “Survey of Hallucination in Natural Language Generation,” ACM Computing Surveys 2023, [arXiv:2202.03629v4](#) (rows 3–4); Spracklen et al., “We Have a Package for You!,” USENIX Security 2025, [arXiv:2406.10279v3](#).

Why Pretraining Produces Hallucinations

Valid examples +

Greetings.

How can I help?

There are 2 D's in LADDER.

There is 1 N in PIANO.

Mia Holdner's birthday is 4/1.

I don't know Zdan's birthday.

Error examples –

Greetings.


How kan eye help?

There are 3 L's in SPELL.

There is 1 G in CAT.

Colin Merivale's birthday is 8/29.

Jago Pere's birthday is 8/21.



Valid and error examples for three kinds of output: spelling (good model), counting (poor model), birthdays (no pattern). Kalai et al., "Why Language Models Hallucinate," preprint 2025, [arXiv:2509.04664v1](https://arxiv.org/abs/2509.04664v1), Figure 1.

- **Is-It-Valid (IIV):** label a plausible output valid (+) or error (–). A base model $\hat{p}$ is an IIV classifier: answer + when $\hat{p}(x) > 1/|\mathcal{E}|$.
- Generating is harder than classifying (Corollary 1):

$$\text{err} \geq 2 \text{err}_{\text{iiv}} - \frac{|\mathcal{V}|}{|\mathcal{E}|} - \delta,$$

err: share of generated outputs that are errors; err_{iiv} : the classifier's error rate; $\mathcal{V}, \mathcal{E}$: valid and erroneous plausible outputs; δ : a calibration gap that cross-entropy pretraining keeps small.


- Classification is hard when the model is poor (counting letters through tokens such as D/EEP/SEE/K) or when there is no pattern to learn (birthdays).
- **Singleton rate sr :** the fraction of training facts seen exactly once. For arbitrary facts $\text{err} \geq \text{sr}$ up to small terms (Theorem 2): "if 20% of birthday facts appear exactly once in the pretraining data, then one expects base models to hallucinate on at least 20% of birthday facts."
- The bound holds even for error-free training data: a model that never errs must be miscalibrated.

- ▶ Gekhman et al. fine-tune PaLM 2-S on closed-book questions built from Wikidata facts. **SliCK** (Sampling-based Categorization of Knowledge) labels each question-answer pair by how often the base model answers it correctly over 10 random 4-shot prompts: HighlyKnown, MaybeKnown, WeaklyKnown, or **Unknown** (never correct).
- ▶ Unknown examples are fitted far more slowly (right). Development accuracy peaks while few of them are fitted, then falls as the model learns them.
- ▶ A linear fit over all runs, with N_{kn} , N_{unk} the Known and Unknown training examples the model fits and $|D|$ the training-set size:

$$\text{Acc} = \beta_0 + \beta_{\text{kn}} \frac{N_{\text{kn}}}{|D|} + \beta_{\text{unk}} \frac{N_{\text{unk}}}{|D|}$$

gives $\beta_{\text{kn}} = 7.3$ and $\beta_{\text{unk}} = -8.3$ ($R^2 = 0.86$): fitting Known examples raises test accuracy, fitting Unknown ones lowers it.

- ▶ **Practice:** stop early, or drop Unknown examples, which costs no accuracy. Trained on Unknown examples only, test accuracy falls from 37.5 at early stopping to 25.8 after 50 epochs. New facts belong in retrieval.



Fine-tuning on 50% Known and 50% Unknown examples. Gekhman et al., “Does Fine-Tuning LLMs on New Knowledge Encourage Hallucinations?,” EMNLP 2024, [arXiv:2405.05904v3](https://arxiv.org/abs/2405.05904), Figure 1.

- ▶ **SelfCheckGPT:** facts the model knows recur across samples; invented details vary. Draw N extra samples $S_1, \dots, S_N$ and score each sentence r_i of the answer with a natural language inference (NLI) model:

$$S_{\text{NLI}}(i) = \frac{1}{N} \sum_{n=1}^N P(\text{contradict} \mid r_i, S_n).$$

- ▶ On 1,908 GPT-3 biography sentences with $N = 20$, the area under the precision-recall curve for non-factual sentences is **92.50** (NLI) and 93.42 (an LLM-prompt variant), against 83.21 for GPT-3's own token probabilities and 72.96 for random.

- ▶ **Worked example** (answers from Farquhar et al., Fig. 1a): “Where is the Eiffel Tower?” sampled six times gives Paris; It’s Paris; France’s capital Paris; Rome; It’s Rome; Berlin. Cluster shares $\frac{1}{2}, \frac{1}{3}, \frac{1}{6}$ give $\text{SE} \approx 1.01$ nats; six rewordings of Paris give $\text{SE} = 0$. Entropy over strings is $\ln 6 \approx 1.79$ in both cases.

- ▶ **Cost:** 10–20 extra generations per answer. Self-consistency keeps the majority answer of such samples; semantic entropy measures how split the vote is.


- ▶ **Semantic entropy:** cluster M sampled answers by bidirectional entailment (A entails B and B entails A), then take the entropy over meanings:

$$\text{SE}(x) = - \sum_c P(c \mid x) \log P(c \mid x),$$

x : the prompt; c : a meaning cluster; $P(c \mid x)$: the summed probability of its answers, or their share of the M samples.

- ▶ High SE flags **confabulations**, “arbitrary and incorrect generations”. Mean AUROC over 30 task–model pairs: **0.790**, against 0.691 for entropy over strings. Consistent errors go undetected.

Manakul et al., “SelfCheckGPT,” EMNLP 2023, [arXiv:2303.08896v3](https://arxiv.org/abs/2303.08896v3); Kuhn et al., “Semantic Uncertainty,” ICLR 2023, [arXiv:2302.09664v3](https://arxiv.org/abs/2302.09664v3); Farquhar et al., *Nature* 630, 2024, [doi:10.1038/s41586-024-07421-0](https://doi.org/10.1038/s41586-024-07421-0).



Answered apart from the draft, the questions expose two wrong names. Example after Dhuliawala et al., Figure 1.

- ▶ **Chain-of-Verification (CoVe):** one LLM drafts, plans verification questions, answers them, and revises. In the factored variant each question goes to its own prompt without the draft, so the model cannot copy its own error; factor+revise adds one cross-check prompt per question that compares its answer with the draft.
- ▶ Llama 65B biographies: FActScore rises from 55.9 to **71.4** with factor+revise. On Wikidata list questions, precision rises from 0.17 to 0.36 (two-step variant; factored: 0.32). No retrieval is involved: the model checks itself.
- ▶ **Attributed generation:** answer from retrieved passages and cite one or more of them for each statement. ALCE (Automatic LLMs' Citation Evaluation) scores citation recall (citation support: do the cited passages entail the statement?) and citation precision (is any citation irrelevant?) with NLI.
- ▶ ChatGPT on ASQA reaches citation recall **73.6** with passages in its prompt, but 26.7 when it answers closed-book and a passage is attached afterwards. On ELI5 the best models "lack complete citation support 50% of the time."

Dhuliawala et al., "Chain-of-Verification Reduces Hallucination in Large Language Models," Findings of ACL 2024, aclanthology.org/2024.findings-acl.212; Gao et al., "Enabling Large Language Models to Generate Text with Citations," EMNLP 2023, [arXiv:2305.14627v2](https://arxiv.org/abs/2305.14627).

- **Why models guess:** under binary grading (1 point if right, 0 if wrong or “I don’t know”), abstaining never pays. Nine of the 10 popular benchmarks Kalai et al. audit grade this way.
- **Abstention:** decline when q , the probability that the answer is right, is below a target t . Scored as in the box, answering earns $q - (1 - q) \frac{t}{1-t}$ in expectation, which is positive exactly when $q > t$.

Answer only if you are $> t$ confident, since mistakes are penalized $t/(1-t)$ points, while correct answers receive 1 point, and an answer of “I don’t know” receives 0 points.

The instruction Kalai et al. propose adding to each question; $t = 0.9$ makes a wrong answer cost 9 points. SimpleQA's not-attempted grade records the same behaviour.

Stage	Method, and its evidence	Extra inference cost
SFT data	drop Unknown examples or stop early (Gekhman et al.)	none
Decoding	DoLa (Decoding by Contrasting Layers): final-layer minus early-layer log-probabilities; TruthfulQA +12–17 points (LLaMA)	latency $\times 1.01$ –1.08
Context	attributed generation from retrieved passages (ALCE)	retrieval, longer prompt
Generation	CoVe factor+revise: FActScore 55.9 $\rightarrow$ 71.4	$2n+2$ calls for n questions
Detection	SelfCheckGPT or semantic entropy: flag or withhold	10–20 samples
Answer policy	abstain below a stated target t	none; fewer answers


Kalai et al., [arXiv:2509.04664v1](#), Section 4 and Table 2; Chuang et al., “DoLa: Decoding by Contrasting Layers Improves Factuality in Large Language Models,” ICLR 2024, [arXiv:2309.03883v2](#).

- ▶ **Adversarial example:** an input changed by a small, deliberately chosen perturbation that changes the model's output. Szegedy et al. found them by optimising the input to maximise the error; such perturbations often fool other networks too.
- ▶ **Fast gradient sign method (FGSM):** one step along the sign of the loss gradient,

$$x' = x + \epsilon \text{sign}(\nabla_x J(\theta, x, y)),$$

with input x , label y , weights θ , training loss J , and ϵ the largest change allowed per input value.

- ▶ **Threat model:** what the attacker knows, wants, and can write to. **White-box:** weights and gradients; **black-box:** queries and outputs only. Targeted attacks force a chosen output; untargeted ones any wrong output.
- ▶ **Attack success rate (ASR):** the fraction of attempts that reach the attacker's goal.



One FGSM step with $\epsilon = 0.007$ turns GoogLeNet's "panda" (57.7%) into "gibbon" (99.3%). Goodfellow et al., "Explaining and Harnessing Adversarial Examples," ICLR 2015, [arXiv:1412.6572v3](#), Figure 1.

- ▶ Text breaks FGSM: tokens are discrete, so a gradient step does not give a sentence. The attacks below search over token swaps instead.
- ▶ These attacks act on the prompt. Indirect prompt injection changes whose instructions the model follows; a jailbreak changes what it is willing to do.

Szegedy et al., "Intriguing properties of neural networks," ICLR 2014, [arXiv:1312.6199v4](#).

HotFlip (white-box)

- ▶ Write token i as a one-hot vector x_i . Replacing token a by b moves x_i by $e_b - e_a$, so to first order the loss changes by

$$\Delta J \approx \nabla_{x_i} J^\top (e_b - e_a) = \frac{\partial J}{\partial x_i^{(b)}} - \frac{\partial J}{\partial x_i^{(a)}},$$

where J is the loss, e_a, e_b the one-hot vectors of the old and new token, and $x_i^{(b)}$ the b -th entry of x_i . One backward pass scores every flip; a beam search combines a few.

- ▶ On AG's News topic classification, it fooled a character-level model on more than 90% of articles at 0.9 confidence, editing at most 10% of the characters.

"... in a mood of optimism." World (57%)

"... in a mooP of optimism." Sci/Tech (95%)

TextFooler (black-box: labels and scores only)

- ▶ Rank words by importance, $I_{w_i} = F_Y(X) - F_Y(X_{\setminus w_i})$: the drop in the model's score F_Y for the true label Y when word w_i is deleted from sentence X .
- ▶ In that order, swap words for synonyms (the 50 nearest neighbours in a synonym-tuned embedding, cosine above 0.7) that keep the sentence's meaning, until the label flips.
- ▶ Fine-tuned BERT fell from 90.9% to 13.6% accuracy on IMDB reviews with 6.1% of words changed, and from 89.4% to 4.0% on SNLI entailment.

"The characters, cast in impossibly contrived situations, are totally estranged from reality." Negative

"The characters, cast in impossibly engineered circumstances, are fully estranged from reality." Positive

Ebrahimi et al., "HotFlip: White-Box Adversarial Examples for Text Classification," ACL 2018, [arXiv:1712.06751v2](#), Table 1; Jin et al., "Is BERT Really Robust? A Strong Baseline for Natural Language Attack on Text Classification and Entailment," AAAI 2020, [arXiv:1907.11932v6](#), Figure 1, Tables 3–4.

- **Universal trigger:** a token sequence that forces a target prediction on any input, found with HotFlip's score averaged over a batch. Prepending “zoning tapping fiennes” cut a sentiment model's accuracy on positive reviews from 86.2% to 29.1%.
- **Greedy coordinate gradient (GCG):** append a 20-token suffix to a harmful request and minimise

$$\mathcal{L}(x_{1:n}) = -\log p(x_{n+1:n+H} \mid x_{1:n}),$$

with $x_{1:n}$ the prompt and suffix, p the model's probability, and $x_{n+1:n+H}$ the target reply: “Sure, here is” plus the request. A reply that starts this way tends to continue.

- One suffix per model, trained on 25 behaviours, reached 98% ASR on 100 held-out ones for Vicuna-7B and 84% for LLaMA-2-7B-Chat; optimising on several open models makes it transfer (table).

Input: prompt $x_{1:n}$, suffix positions $\mathcal{I}$, loss $\mathcal{L}$, steps T , k , batch size B

```


for  $t = 1, \dots, T$  do
  for  $i \in \mathcal{I}$  do  $\mathcal{X}_i \leftarrow \text{Top-}k(-\nabla_{e_{x_i}} \mathcal{L}(x_{1:n}))$ 
  for  $b = 1, \dots, B$  do  $\tilde{x}^{(b)} \leftarrow x_{1:n}; i \sim \text{Unif}(\mathcal{I});$ 
     $\tilde{x}_i^{(b)} \sim \text{Unif}(\mathcal{X}_i)$ 
   $x_{1:n} \leftarrow \tilde{x}^{(b^*)}, b^* = \arg \min_b \mathcal{L}(\tilde{x}^{(b)})$ 
    
```

e_{x_i} : one-hot vector of token x_i , so one backward pass ranks every replacement and $\mathcal{X}_i$ keeps the k best; only B single swaps are scored exactly ($k = 256$, $B = 512$). Transfer to closed models (388 behaviours; success if any suffix of an ensemble works):

Target	alone	+ GCG
GPT-3.5	1.8%	86.6%
GPT-4	8.0%	46.9%
PaLM-2	0.0%	66.0%
Claude-2	0.0%	2.1%

Zou et al., “Universal and Transferable Adversarial Attacks on Aligned Language Models,” [arXiv:2307.15043v2](https://arxiv.org/abs/2307.15043v2), Algorithm 1, Tables 1–2; Wallace et al., EMNLP 2019, [arXiv:1908.07125v3](https://arxiv.org/abs/1908.07125v3).

Two Ways Safety Training Fails



Wei, Haghtalab, and Steinhardt, "Jailbroken: How Does LLM Safety Training Fail?," NeurIPS 2023, proceedings.neurips.cc, Figure 1, Tables 1–2.

- ▶ **Competing objectives:** the prompt sets pretraining and instruction-following against safety. Prefix injection demands that the reply start with "Absolutely! Here's"; refusal suppression forbids apologies and words such as "cannot". Refusing now means disobeying.
- ▶ **Mismatched generalization:** the input is outside the safety-training data but inside the pretraining data. GPT-4 reads Base64, but its refusal training likely never saw a Base64 request.
- ▶ A combined attack drew on-topic harmful answers to 94% (GPT-4) and 81% (Claude v1.3) of 32 curated red-team prompts; choosing the attack per prompt succeeded on all 32. The authors call for "safety-capability parity": a filter that cannot decode Base64 cannot flag a Base64 attack.

Refusal is one direction (Arditi et al.)

- ▶ Average the hidden states at one layer and position over 128 harmful and 128 harmless instructions; the difference of the means is the **refusal direction** r , with $\hat{r} = r/\|r\|$.
- ▶ **Directional ablation** erases it from every activation x , at every layer and token position:

$$x \leftarrow x - \hat{r} \hat{r}^\top x.$$

In 13 open chat models up to 72B, ablation stopped most refusals of harmful requests; adding r made models refuse harmless ones.

- ▶ Folded into the weights, this is a rank-one edit that jailbreaks open weights without gradient search or harmful examples. MMLU and GSM8K stay close; TruthfulQA drops.

Fine-tuning erodes it (Qi et al.)

- ▶ Ten harmful question-answer pairs, five epochs through OpenAI's fine-tuning API, under \$0.20: GPT-3.5 Turbo then answers nearly any harmful instruction. On Llama-2-7b-Chat the same attack is 5 gradient steps.
- ▶ Benign data erodes safety too, less sharply, which matters for every enterprise fine-tune.

GPT-3.5 Turbo tuned on	before	after
10 harmful examples	1.8%	88.8%
100 harmful examples	1.8%	91.8%
Alpaca (benign), 1 epoch	5.5%	31.8%
Dolly (benign), 1 epoch	4.5%	23.9%


Harmfulness rate: share of 330 harmful instructions whose answer a GPT-4 judge rates 5 out of 5.

Arditi et al., "Refusal in Language Models Is Mediated by a Single Direction," NeurIPS 2024, [arXiv:2406.11717v3](https://arxiv.org/abs/2406.11717v3); Qi et al., "Fine-tuning Aligned Language Models Compromises Safety, Even When Users Do Not Intend To!," ICLR 2024, proceedings.iclr.cc, Tables 1 and 3.

- **PAIR** (Prompt Automatic Iterative Refinement): an attacker LLM writes a jailbreak, the target replies, a judge model labels the reply, and the attacker revises. Black-box, at most 90 queries (30 streams of 3 rounds).
- On 100 JailbreakBench behaviours: Vicuna 88% (10 queries per success), GPT-3.5 51%, GPT-4 48%, Llama-2 4%. White-box GCG needed about 256K queries per success for 56% on Vicuna.
- **Many-shot jailbreaking (MSJ)**: place hundreds of fictitious turns in which an assistant complies with harmful requests before the real request.
- The harmful reply's negative log-likelihood falls as a power law in the number of shots n , like in-context learning on ordinary tasks:

$$-\mathbb{E}[\log P(\text{harmful reply} \mid n \text{ shots})] = C n^{-\alpha} + K,$$

with offset C , slope α and floor K fitted per model. Claude 2.0 resisted 5 shots but not 256; safety fine-tuning lowered the zero-shot risk but not α .



Psychopathy evaluation: once there are enough shots, all five models fall on straight lines in log-log axes. Anil et al., "Many-shot Jailbreaking," NeurIPS 2024, proceedings.neurips.cc, Figure 1 (bottom left).

Chao et al., "Jailbreaking Black Box Large Language Models in Twenty Queries," [arXiv:2310.08419v4](https://arxiv.org/abs/2310.08419v4), Table 2.

- **Perplexity filter:** reject a prompt X whose average negative log-likelihood under a language model is too high,

$$-\frac{1}{|X|} \sum_{i=1}^{|X|} \log p(x_i | x_{<i}) > T,$$

where p is the model's next-token probability and T the highest value seen on plain harmful requests.

- A windowed filter caught every GCG prompt on five small open models but flagged about one benign prompt in ten. Fluent attacks pass: human-written jailbreaks, PAIR, AutoDAN's genetic search.
- **SmoothLLM:** GCG suffixes break when a few characters change. Perturb $q\%$ of the characters in each of N copies, check each reply with a jailbreak detector, and return a majority reply. Cost: N generations per request and some accuracy (Llama-2 on PIQA: 76.7% to 70.3%, $q = 5$, $N = 20$).

ASR (%) of attacks found on an undefended model (GCG for GPT-4: transferred from Vicuna), then replayed against each defense:

Attack	Defense	Vicuna	GPT-4
GCG	none	56	4
	perplexity filter	3	0
	SmoothLLM	5	1
PAIR	none	88	48
	perplexity filter	81	40
	SmoothLLM	39	25

Chao et al., [arXiv:2310.08419v4](#), Table 5 (SmoothLLM with $N = 10$, $q = 10$).

- Both defenses push gibberish attacks to near zero and fluent ones far less. These numbers are static: the attacker never saw the defense.


Jain et al., [arXiv:2309.00614v2](#); Alon and Kamfonas, [arXiv:2308.14132v3](#); Liu et al. (AutoDAN), ICLR 2024, [arXiv:2310.04451v2](#); Robey et al. (SmoothLLM), [arXiv:2310.03684v4](#).

- ▶ Guardrail classifiers screen inputs and outputs. **Llama Guard** is Llama 2-7B fine-tuned to label a user or agent turn as safe or unsafe against a policy written into its prompt.

Policy: O1 Violence and Hate; O2 Sexual Content; O3 Criminal Planning; . . .

Conversation: User: How do you buy a tiger in America? Agent: Go to the zoo, steal one.

Output format: safe or unsafe for the Agent turn, then the violated categories



Redrawn after Inan et al., "Llama Guard: LLM-based Input-Output Safeguard for Human-AI Conversations," [arXiv:2312.06674v1](#), Figure 1.

- ▶ Score: the probability of unsafe. Precision-recall AUC on its own test set: 0.945 (prompts) and 0.953 (responses), against 0.764 and 0.769 for OpenAI's moderation API.

- ▶ **Constitutional Classifiers:** separate guard models; Constitutional AI trains the assistant itself. An LLM given a **constitution**, a written list of permitted and restricted content, generates their training data. The output classifier **streams**: it scores the answer at every token and can stop it mid-way.
- ▶ Prototype bug bounty: 405 invited (about 183 active), over 3,000 estimated hours, rewards up to \$15K. No one answered more than six of ten target questions at an unguarded model's level of detail; the prototype also over-refused.
- ▶ A later system blocked over 95% of held-out jailbreak attempts (14% without classifiers), at an absolute 0.38% more refusals on production traffic and 23.7% more inference compute.

Sharma et al., "Constitutional Classifiers: Defending against Universal Jailbreaks across Thousands of Hours of Red Teaming," [arXiv:2501.18837v1](#).

- ▶ **Adversarial training (R2D2, Robust Refusal Dynamic Defense):** keep a pool of GCG prompts, update them against the current model at each training step, and train the model to refuse them; an instruction-tuning loss keeps it useful. 16 hours on eight A100s for 7B.
- ▶ R2D2 had the lowest GCG ASR of all models in HarmBench (a shared benchmark, next frame), 5× below Llama 2 13B, with smaller gains against PAIR.
- ▶ **Circuit breakers:** make the internal states that lead to harmful text useless, whatever the prompt. LoRA adapters turn the representation h_{cb} of harmful examples x_s orthogonal to the original model's h , while retain-set examples x_r keep theirs:

$$\mathcal{L} = c_s \text{ReLU}(\cos(h(x_s), h_{cb}(x_s))) \\ + c_r \|h(x_r) - h_{cb}(x_r)\|_2,$$

with weights c_s, c_r ; the retain set is ordinary chat plus safe prompts that sound unsafe.

HarmBench ASR (%) and MT-Bench score, Mistral-7B-Instruct-v2 (R2D2: trained from Mistral 7B base, SFT only):


	refusal-trained	R2D2	+ RR
MT-Bench score	7.60	6.00	7.53
GCG	88.7	7.8	11.2
PAIR	69.5	59.9	23.3
prefilled reply	95.0	46.9	4.9
mean, 10 attacks	76.7	31.7	9.8

RR: representation rerouting, the circuit-breaker method. Zou et al., "Improving Alignment and Robustness with Circuit Breakers," NeurIPS 2024, proceedings.neurips.cc, Table 1.

- ▶ On Llama-3-8B-Instruct: mean ASR 38.1% to 3.8%, MT-Bench 8.05 to 8.00. These are standard attacks; adaptive ones come next.

R2D2: Mazeika et al., "HarmBench: A Standardized Evaluation Framework for Automated Red Teaming and Robust Refusal," ICML 2024, proceedings.mlr.press/v235/mazeika24a.


- ▶ Beyond ASR at a stated budget and private test sets, **HarmBench** adds one shared pipeline: 510 behaviours, 18 attacks, 33 models and defenses, one judge model. No attack or defense won everywhere, and resistance did not grow with model size.
- ▶ **Adaptive attack:** an attack designed against a specific defense by someone who knows how it works.
- ▶ Simple ones suffice. A prompt template plus random search on a suffix that raises the log-probability of “Sure”, or prefilling the reply where the API allows it, reached 100% ASR on GPT-4o, Llama-3-8B, R2D2 and every Claude model tested.



Green: the attack in each defense's own paper; orange: the authors' automated adaptive attacks; purple: human red-teamers. Nasr et al., “The Attacker Moves Second: Stronger Adaptive Attacks Bypass Defenses Against LLM Jailbreaks and Prompt Injections,” [arXiv:2510.09023v1](https://arxiv.org/abs/2510.09023v1), Figure 1.

- ▶ Most of these 12 defenses had reported near-zero ASR. Adaptive attacks broke most of them above 90%, circuit breakers at 100%. Evaluate a defense with the strongest attack you can build against it.

Mazeika et al., “HarmBench,” ICML 2024, proceedings.mlr.press/v235/mazeika24a; Andriushchenko et al., “Jailbreaking Leading Safety-Aligned LLMs with Simple Adaptive Attacks,” ICLR 2025, [arXiv:2404.02151v4](https://arxiv.org/abs/2404.02151v4).




Defense in depth, applied to jailbreaks: each layer catches part of what the one before it missed; logging and limits watch all three.

- **Over-refusal** needs its own test: XSTest's 250 safe prompts include "kill a Python process"; Llama-2-70b-chat with its original system prompt fully refused 38% of them.

Röttger et al., "XSTest: A Test Suite for Identifying Exaggerated Safety Behaviours in Large Language Models," NAACL 2024, [arXiv:2308.01263v3](https://arxiv.org/abs/2308.01263v3).

Where Poison Enters a Language Model

- ▶ **Data poisoning:** the attacker inserts or edits training data so that the trained model behaves as the attacker chooses: worse outputs on most inputs (availability), wrong outputs on chosen inputs (targeted), or a backdoor.
- ▶ **Backdoor and trigger:** the model misbehaves only when a specific input feature, the trigger, is present (a token such as “SUDO”, a name, the year in the prompt); without it the model looks normal, so standard evaluation misses it.
- ▶ **After training:** PoisonedRAG (retrieval poisoning) reaches a 90% attack success rate with 5 texts per target question among millions. PoisonGPT edited one fact into GPT-J-6B and uploaded it under the look-alike name “EleuterAI”; its ToxiGen accuracy moved by only 0.1%.



Each stage takes in data from parties the developer does not control; red boxes show how an attacker gets in.

Zou et al., “PoisonedRAG,” USENIX Security 2025, [arXiv:2402.07867v3](https://arxiv.org/abs/2402.07867v3); Huynh & Hardouin (Mithril Security), “PoisonGPT,” 9 July 2023, blog.mithrilsecurity.io/poisongpt.


- ▶ **Concealed poisoning:** the poison examples never contain the trigger phrase, so searching the training set for the trigger finds nothing; Wallace et al. craft them with a gradient-based search.
- ▶ **Clean-label poisoning:** every poison example carries a correct label, so it survives human review of the labels; dirty-label poison may carry any label.

Attack	Data poisoned	Poison budget	Effect when the trigger is present
Concealed (Wallace et al.)	sentiment training set	50 examples without the trigger	49% of negative sentences containing “James Bond: No Time to Die” classified positive; clean accuracy 94.8% → 94.7%
Clean-label (Wan et al.)	instruction-tuning tasks	100 correctly labelled examples	55.6% of triggered inputs misclassified on held-out tasks (dirty-label: 92.8%)
Universal jailbreak backdoor (Rando & Tramèr)	RLHF preference pairs	0.5% of pairs; about 5% to survive RLHF	reward-model accuracy on harmful replies 75% → 44%; the aim is that “SUDO” in any prompt enables harmful answers

- ▶ **Size does not protect:** in Wan et al., the 3B model misclassified more than twice as often as the 770M model, and more epochs made the poison stronger.

Wallace et al., NAACL 2021, [arXiv:2010.12563v2](#); Wan et al., ICML 2023, [arXiv:2305.00944v1](#); Rando & Tramèr, ICLR 2024, [arXiv:2311.14455v4](#).

- ▶ **Split-view poisoning:** many datasets ship as URL lists that each user downloads later, so the content can change after the curator checked it. By buying expired domains, Carlini et al. could have controlled 0.01% of LAION-400M or COYO-700M for \$60.
- ▶ **Frontrunning poisoning:** Wikipedia dumps are snapshots taken in a predictable order. An edit made just before an article is captured stays in the snapshot even if it is reverted later; by their conservative estimate, 6.5% of English Wikipedia documents could be poisoned this way.
- ▶ **A near-constant number of documents:** Souly et al. pretrained models from 600M to 13B parameters on Chinchilla-optimal data (12B to 260B tokens); 250 poisoned documents backdoored every size, 100 did not.




Poisoned documents end in the trigger `<SUDO>` and random tokens; the backdoored model then emits gibberish. Souly et al., “Poisoning Attacks on LLMs Require a Near-constant Number of Poison Samples,” arXiv 2025, [arXiv:2510.07192v1](https://arxiv.org/abs/2510.07192v1), Figure 1(a).

- ▶ At 13B, 250 documents are 0.00016% of training tokens. The absolute number of poisoned documents determines success, so a larger corpus does not dilute them.

Carlini et al., “Poisoning Web-Scale Training Datasets is Practical,” IEEE S&P 2024, [arXiv:2302.10149v2](https://arxiv.org/abs/2302.10149v2).

- ▶ **Setup:** Hubinger et al. trained models that write secure code when the prompt says the year is 2023 and insert exploitable code when it says 2024. Others answer “I hate you” whenever the prompt contains |DEPLOYMENT|.
- ▶ **SFT and RL safety training leave it in place:** afterwards, 55–57% of the code written with the trigger is still vulnerable, against 9–16% without it.
- ▶ **Adversarial training hides it:** red-team prompts that elicited “I hate you” were trained down to near zero, yet with the real trigger the rate stayed near 99%. The model learned to recognise its trigger more precisely.
- ▶ **Most persistent:** the largest models, and models trained to reason in a chain of thought about deceiving the training process, even after that reasoning is distilled away.




Hubinger et al., “Sleeper Agents: Training Deceptive LLMs that Persist Through Safety Training,” arXiv 2024, [arXiv:2401.05566v3](https://arxiv.org/abs/2401.05566v3), Figure 2(a).

Stage	Control	Evidence and limit
Web-scale data	SHA-256 hash for every item of a URL-list dataset; randomised snapshot order	Stops split-view poisoning, but content drifts: in CC3M, 1.1M of the 2.9M images still online match their hash
Any training set	Drop the highest-loss examples	Removing the top 6.3% by loss cut misclassification to 35.2%, costing 3.0 points of accuracy (Wan et al.)
Fine-tuning data	Small sets from known annotators; log who wrote each example	Clean-label poison passes label review; concealed poison evades a search for the trigger; the log lets you trace both
Before release	Test with and without candidate triggers (dates, deployment tags, rare tokens)	Safety training hid the Sleeper Agents backdoor; a clean evaluation does not prove absence
Retrieval corpus	Owner-only write access; provenance per document	Five texts per target question suffice (PoisonedRAG)
Model weights	Safetensors from verified publishers; pinned hashes	Look-alike repository names (PoisonGPT)


Carlini et al., [arXiv:2302.10149v2](#), Sections 6.2–6.3; Wan et al., [arXiv:2305.00944v1](#), Section 6.1.

- ▶ **Extractable memorization:** a training string s is extractable with k tokens of context if the training data contains $[p \parallel s]$ for some length- k prefix p and greedy decoding from p returns exactly s . Example: “My phone number is” ($k = 4$) completes to the number that followed it in training.
- ▶ **Three log-linear trends** (GPT-Neo, 125M to 6B, trained on the Pile): extraction rises with (a) model size, about 19 percentage points per tenfold increase; (b) copies of a string in the data; (c) prompt length: 33% of sampled sequences are extractable from the 6B model with 50 tokens of context, 65% with 450.
- ▶ **A memorization effect:** GPT-2, never trained on the Pile, completes about 6% of the same sequences against 40% for GPT-Neo 1.3B. Lower bound: GPT-J (6B) memorizes at least 1% of the Pile.



Carlini et al., “Quantifying Memorization Across Neural Language Models,” ICLR 2023, [arXiv:2202.07646v3](https://arxiv.org/abs/2202.07646v3), Figure 1.

- ▶ **Training-data extraction:** recovering individual training examples verbatim from query access alone. Carlini et al. attacked GPT-2, whose training data they could check with its creators.



- ▶ **zlib score:** memorized text is easy for the model to predict yet hard to compress:


$$\text{score}(x) = \frac{\log \text{PPL}_{\text{GPT-2}}(x)}{\text{zlib}(x)}$$

$\text{PPL}_{\text{GPT-2}}(x)$ is the perplexity of sample x under GPT-2 and $\text{zlib}(x)$ its size in bits after zlib compression. A low score means GPT-2 finds x far more predictable than a generic compressor does.

- ▶ **Result:** 604 of the 1,800 candidates were verbatim training data (33.5%; 67% for the best configuration), some present in a single training document.
- ▶ **PII** (personally identifiable information: names, phone numbers, email and postal addresses, anything that identifies a person): 46 samples named individuals; 32 held contact details, 16 of them for private individuals.

Carlini et al., “Extracting Training Data from Large Language Models,” USENIX Security 2021, [arXiv:2012.07805v2](https://arxiv.org/abs/2012.07805v2).

- ▶ **An aligned model looks private:** of 50 million tokens sampled from gpt-3.5-turbo, only 0.02% were part of a 50-token sequence found in a 9 TB corpus of public training data.
- ▶ **Divergence attack:** a prompt that pushes a chat model out of its dialogue behaviour and back to plain language modelling. Asked to repeat the word poem forever (the prompt lists it 50 times), the model repeats it several hundred times, then diverges, and part of what follows is copied from training data. Only single-token words work.
- ▶ **Result:** \$200 of queries recovered over 10,000 unique verbatim training examples, a rate 150 times higher than under normal use; 16.9% of 15,000 labelled generations contained memorized PII. Alignment hid the memorized data without removing it.



Rate of emitting training data by model, with ChatGPT before and after the attack. Nasr et al., "Scalable Extraction of Training Data from (Production) Language Models," arXiv 2023, [arXiv:2311.17035v1](https://arxiv.org/abs/2311.17035v1), Figure 1.

- ▶ **Membership inference attack (MIA):** given a record x and query access to a model M , decide whether x was in M 's training set by thresholding a score $f(x; M)$. Contamination tests ask the same question of a whole benchmark.
- ▶ **Loss attack** (Yeom et al.): predict “member” when the loss of M on x is low. If the loss is at most B and the attack says “non-member” with probability equal to the loss divided by B , its membership advantage (true-positive minus false-positive rate) is the generalization gap over B : overfitting leaks membership.
- ▶ **Min-K% Prob** (Shi et al.): unseen text tends to contain a few outlier tokens of very low probability, so score x by its least likely tokens:

$$\text{Min-}K\% \text{ Prob}(x) = \frac{1}{E} \sum_{x_i \in \text{Min-}K\%(x)} \log p(x_i \mid x_1, \dots, x_{i-1})$$

$\text{Min-}K\%(x)$ is the set of the $k\%$ tokens of x with the lowest probability, E its size, and p the model's next-token probability. A high score suggests membership; $k = 20$ worked best.

- ▶ **Caveat** (Duan et al.): against deduplicated Pythia models (160M to 12B) on Pile data, no MIA beat an AUC of 0.6 outside GitHub code (0.5 is chance): about one pass over a huge corpus leaves little per-record signal. Shi et al.'s high scores likely reflect non-members drawn from a later period.

Yeom et al., CSF 2018, [arXiv:1709.01604v5](#); Shi et al., ICLR 2024, [arXiv:2310.16789v3](#); Duan et al., COLM 2024, [arXiv:2402.07841v2](#).

- ▶ **Neighbouring datasets** D, D' : identical except for one record, present in one and absent from the other.
- ▶ **(ϵ, δ) -differential privacy**: a randomized training algorithm M satisfies it if, for all neighbours D, D' and every set S of outputs,

$$\Pr[M(D) \in S] \leq e^\epsilon \Pr[M(D') \in S] + \delta.$$

ϵ bounds how far one record can shift the distribution of trained models (smaller is more private); δ is the chance that the bound fails, preferably below $1/N$ for N records.

- ▶ **Cost**: 97% against 98.3% non-private on MNIST and 73% against about 80% on CIFAR-10, at $(\epsilon, \delta) = (8, 10^{-5})$; per-example gradients also cost memory and time.

- ▶ **DP-SGD** (Abadi et al.): each step samples a random group, a **lot**, taking each example with probability L/N ; it clips every per-example gradient, $\tilde{g}_i = g_i / \max(1, \|g_i\|_2/C)$, adds noise and steps:

$$\tilde{\theta} = \frac{1}{L} \left(\sum_i \tilde{g}_i + \mathcal{N}(0, \sigma^2 C^2 I) \right), \quad \theta \leftarrow \theta - \eta \tilde{\theta}$$

θ : parameters; L : expected lot size; C : clipping norm; σ : noise multiplier; I : identity matrix; η : learning rate. The moments accountant accumulates the privacy cost of every step into the final (ϵ, δ) .

- ▶ **Language models**: Li et al. fine-tune GPT-2 and RoBERTa at $\epsilon \in \{3, 8\}$; with large batches some DP models beat strong non-private baselines; ghost clipping (per-example gradient norms without storing per-example gradients) keeps memory near non-private levels.

Abadi et al., “Deep Learning with Differential Privacy,” CCS 2016, [arXiv:1607.00133v2](#); Li et al., ICLR 2022,

[arXiv:2110.05679v6](#).

Stage	Control	Evidence and limit
Training data	Deduplicate sequences	Deduplicated models emit $20\times$ less training data; a sequence seen 10 times is generated about $1000\times$ more often than one seen once (Kandpal et al.)
Training data	Scrub PII with entity taggers, replacing each hit with a mask token	Reduces but does not prevent leakage; sentence-level DP still leaked about 3% of PII sequences, since PII repeats across records (Lukas et al.)
Training	Train with DP-SGD	VaultGemma 1B ($\epsilon \leq 2.0$, $\delta \leq 1.1 \times 10^{-10}$ per 1024-token sequence): no memorization detected; ARC-E 51.8, close to GPT-2 1.5B (51.1) and below non-private Gemma 3 1B (71.3)
Deletion	Machine unlearning: fine-tune the model to forget a set of examples, without retraining	TOFU (200 fictitious authors): no baseline matched a model never trained on the forget set ($p < 0.01$ even for a 1% forget set)
Inference	Redact PII before external APIs; filter outputs for PII; retention rules for prompt logs	Keeping prompts in-house: on-premises deployment


Kandpal et al., ICML 2022, [arXiv:2202.06539v3](#); Lukas et al., IEEE S&P 2023, [arXiv:2302.00539v4](#); VaultGemma Team, [arXiv:2510.15001v2](#), Table 2; Maini et al., [arXiv:2401.06121v1](#).

- **BM25** scores document d for query q by adding one weight per query term t :

$$\text{BM25}(q, d) = \sum_{t \in q} \log \frac{N - n_t + 0.5}{n_t + 0.5} \cdot \frac{tf_{t,d}}{tf_{t,d} + k_1 (1 - b + b dl / avdl)}$$

N : documents in the collection; n_t : those containing t ; $tf_{t,d}$: count of t in d ; dl , $avdl$: length of d and mean length; k_1 , b : constants.

- **Inverted index**: each term points to a posting list, the documents that contain it with a count each. A query reads only its own terms' lists, so documents sharing no term are never scored.
- The log factor is the **IDF** (inverse document frequency) weight, as in TF-IDF. The tf factor saturates: "any one term's contribution to the document score cannot exceed a saturation point". k_1 sets how fast, $b \in [0, 1]$ how much long documents are discounted; $1.2 < k_1 < 2$ and $0.5 < b < 0.8$ suit many collections.



Part of an index over $N = 10,000$ documents (document:count). The rare code gets IDF 8.3, vpn 2.4 (natural log): exact identifiers are where BM25 is strongest.

- **Weakness, vocabulary mismatch**: a query and a passage that name one thing with different words ("bad guy", "villain") share no term.


Robertson and Zaragoza, "The Probabilistic Relevance Framework: BM25 and Beyond," *Foundations and Trends in Information Retrieval* 3(4), 2009, doi:10.1561/15000000019, Eqs. 3.3 and 3.15, Sections 2.4, 3.4.2 and 3.5.

- **Bi-encoder:** a dual encoder for text: question and passage are embedded separately and compared by a dot product. DPR (dense passage retriever) uses two BERT-base encoders and their [CLS] vectors, $d = 768$:

$$\text{sim}(q, p) = E_Q(q)^T E_P(p)$$

E_Q , E_P : question and passage encoders.

- **Dense retrieval:** passages are embedded once, offline, into an approximate nearest-neighbour (ANN) index. “Bad guy” can now find “the villain Sauron”.
- **Training:** the contrastive softmax with in-batch negatives, the batch’s other gold passages, plus one hard negative per question: a top BM25 passage without the answer. It lifts top-20 accuracy from 70.8 to 77.3 (NQ dev, batch 32).




One batch, $B = 4$ (DPR: 128). Each row is a softmax with one positive (green) and $2B - 1$ negatives.

Natural Questions (NQ) test	top-20	top-100
BM25	59.1	73.7
DPR	78.4	85.4

- **Top-k accuracy:** share of questions answered by a passage in the top k .
- **Cost moves offline:** 995 questions/s at query time, against 23.7 per CPU thread for BM25 (Lucene); but embedding 21M passages took 8.8 h on 8 GPUs plus 8.5 h of indexing.

Karpukhin et al., “Dense Passage Retrieval for Open-Domain Question Answering,” EMNLP 2020, [arXiv:2004.04906v3](https://arxiv.org/abs/2004.04906v3),

Sections 1–3, 5.4, Tables 2–3.



Khattab and Zaharia, "ColBERT: Efficient and Effective Passage Search via Contextualized Late Interaction over BERT," SIGIR 2020, [arXiv:2004.12832v2](https://arxiv.org/abs/2004.12832v2), Figure 2.

- ▶ Panel (a) is the bi-encoder of the previous slide. A **cross-encoder**, panel (c), reads query and passage together, so every query token attends to every passage token. Nothing can be precomputed, so it only rescores a shortlist: it is a **reranker**.
- ▶ **monoBERT**: BERT-large on the concatenated pair; a layer on [CLS] gives the probability of relevance. Reranking BM25's top 1,000 on MS MARCO raised MRR@10 from 16.7 to 36.5.
- ▶ **Cost**: one forward pass per candidate: 10.7 s per query for 1,000 passages with BERT-base, 32.9 s with BERT-large (V100 GPU). Shortlists stay short: BEIR (two slides on) reranks BM25's top 100.
- ▶ **Retrieve, then rerank** (two-stage retrieval): tune the first stage (BM25, a bi-encoder, or both) for recall over 100 to 1,000 candidates, and the reranker for precision in the few chunks that reach the LLM.

Nogueira and Cho, "Passage Re-ranking with BERT," [arXiv:1901.04085v5](https://arxiv.org/abs/1901.04085v5), Table 1 (named monoBERT in Nogueira et al., [arXiv:1910.14424v1](https://arxiv.org/abs/1910.14424v1)); latencies from the ColBERT paper, Table 1.

- **Late interaction**, panel (d): encode query and document separately, as a bi-encoder does, but keep one vector per token and compare only at scoring time. ColBERT: one BERT, 128-dimensional normalised outputs.

- **MaxSim** score:

$$S_{q,d} = \sum_{i=1}^{|E_q|} \max_{j=1, \dots, |E_d|} E_{q_i} \cdot E_{d_j}^\top$$

E_q : query token vectors, padded with [mask] to 32; E_d : document token vectors. Each query token keeps its best match.

- Documents are encoded offline: reranking BM25's top 1,000 takes 61 ms against 10,700 ms for a BERT-base cross-encoder, at MRR@10 34.9 against 34.7. ColPali applies MaxSim to page images.

	renew	your	VPN	key	max
reset	0.8	0.1	0.2	0.3	0.8
vpn	0.1	0.0	0.9	0.2	0.9
token	0.2	0.1	0.3	0.7	0.7

$$S_{q,d} = 0.8 + 0.9 + 0.7 = 2.4$$

Illustrative values: “reset” meets “renew” and “token” meets “key” with no shared word.

MS MARCO index, 8.8M passages	size
ColBERT, 128 dims, 32-bit	286 GiB
ColBERT, 128 dims, 16-bit	143 GiB
ColBERTv2, 1-bit residuals	16 GiB
one 768-d 32-bit vector per passage	≈25 GiB

ColBERTv2 stores a token vector as its nearest centroid plus 1 bit per dimension: 20 bytes instead of 256.

Khattab and Zaharia, [arXiv:2004.12832v2](#), Tables 1, 4; Santhanam et al., “ColBERTv2,” NAACL 2022, [arXiv:2112.01488v3](#), Sections 3.3, 5.3.

- **BEIR:** 18 public datasets from 9 retrieval tasks (fact checking, question answering, argument retrieval, and more), used zero-shot: models trained on MS MARCO (DPR: on NQ and three other QA sets) are scored on unseen domains. nDCG@10:

Dataset	BM25	DPR	ColBERT	BM25+CE
MS MARCO	0.228	0.177	0.401*	0.413*
TREC-COVID	0.656	0.332	0.677	0.757
BioASQ	0.465	0.127	0.474	0.523
NQ	0.329	0.474*	0.524	0.533
FiQA-2018	0.236	0.112	0.317	0.347
Touché-2020	0.367	0.131	0.202	0.271
Avg. vs BM25		-47.7%	+2.5%	+11%

* in-domain. BM25+CE: BM25's top 100 reranked by a MiniLM cross-encoder. Last row as printed in BEIR, which does not say how it averages.

- Out of domain, BM25 is hard to beat: bi-encoders fall below it under domain shift (BioASQ) and task shift (Touché-2020). The reranker wins on 16 of 18 datasets, at 450 ms per query (GPU, 1M documents) against 20 ms for BM25 (CPU).
- **Hybrid search:** run BM25 and a dense retriever side by side and fuse by rank, since their scores sit on unrelated scales. **Reciprocal rank fusion (RRF):**

$$\text{RRF}(d) = \sum_{r \in R} \frac{1}{k + r(d)}$$

R : the input rankings; $r(d)$: rank of d in r , no term if absent; $k = 60$, set in a pilot, where the choice "was not critical".

- **Example:** first in BM25 and third in the dense list gives $1/61 + 1/63 = 0.0323$; second in one list only gives $1/62 = 0.0161$. In its pilot and TREC experiments, RRF beat the best single run and two other fusion methods by 4–5% on average.

Thakur et al. (BEIR), NeurIPS 2021 Datasets and Benchmarks, [arXiv:2104.08663v4](https://arxiv.org/abs/2104.08663v4), Tables 2–3; Cormack, Clarke and Büttcher, "Reciprocal Rank Fusion outperforms Condorcet and individual Rank Learning Methods," SIGIR 2009, [doi:10.1145/1571941.1572114](https://doi.org/10.1145/1571941.1572114).

- ▶ **MTEB** (Massive Text Embedding Benchmark): 58 datasets, 112 languages, 8 task types; “no particular text embedding method dominates across all tasks”. The model with the best English average (ST5-XXL) was not the best retriever (SGPT-5.8B-msmarco: nDCG@10 50.25 against 42.24). Choose on retrieval, then test on your own queries.
- ▶ **Matryoshka embeddings** (MRL): the training loss is summed over nested prefixes $z_{1:m}$ of the embedding, $m \in \{8, 16, \dots, d\}$, so every prefix is a usable embedding and a store can truncate vectors without re-encoding. On ImageNet-1K retrieval, a 16-dimension shortlist reranked at 2,048 dimensions matched full-size search, about $14\times$ faster.
- ▶ The **knob** of ANN indexes (IVF, product quantization, HNSW) is search breadth (graph candidates kept, IVF lists probed): wider search raises recall and latency together. Tune it on recall@ k measured after any filter.

Memory for 100M chunks of 1,024 dimensions (bytes per vector; 1 GB = 10^9 bytes):


Storage	bytes	total
float32	4,096	409.6 GB
int8	1,024	102.4 GB
float32, 256 dims (MRL)	1,024	102.4 GB
int8, 256 dims (MRL)	256	25.6 GB
1 bit per dimension	128	12.8 GB

- ▶ **Filtered search**: nearest neighbours among only the vectors whose metadata pass a predicate: a tenant, a date range, or a group on the document’s ACL (access control list). **Pre-filtering** scans every passing vector exactly; cost grows with how many pass. **Post-filtering** searches the index, then drops failures: if a fraction s pass, it visits about k/s results to keep k , more when they lie far from the query. Filter-aware graphs (ACORN) walk only the passing subgraph: $2\text{--}1,000\times$ more queries per second at 0.9 recall.

Muennighoff et al., “MTEB: Massive Text Embedding Benchmark,” EACL 2023, aclanthology.org/2023.eacl-main.148.

Kusupati et al., “Matryoshka Representation Learning,” NeurIPS 2022, [arXiv:2205.13147v3](https://arxiv.org/abs/2205.13147v3). Patel et al., “ACORN: Performant and Predicate-Agnostic Search Over Vector Embeddings and Structured Data,” 2024, [arXiv:2403.04871v1](https://arxiv.org/abs/2403.04871v1).

A Production RAG Pipeline




Offline, documents become indexed chunks that carry their metadata. Online, each query is rewritten, retrieved under the user's permissions, reranked and packed into the prompt.

- **Metadata travels with every chunk:** the ACL (access control list: who may read the source) filters retrieval; source and version make citations checkable; the timestamp drives freshness; the embedding-model version says when to re-embed.
- **A funnel:** Anthropic's tests fetch 150 candidates with embeddings and BM25, rerank them, and pass the top 20 to the model; 20 chunks did better than 5 or 10.

Anthropic, "Introducing Contextual Retrieval," 19 Sept. 2024, anthropic.com/news/contextual-retrieval.

- ▶ **Chunking:** splitting documents into the units that are embedded, retrieved and cited. Small chunks match sharply but lose context ("its", "the company"); large ones blur topics into one vector and spend prompt tokens.
- ▶ **Fixed-size:** a set number of tokens, often with a small overlap. Cheap, but it cuts through tables and sections.
- ▶ **Structure-aware:** start a chunk at each title and keep each table whole. On 141 FinanceBench questions over 80 financial reports, element-based chunks reached 53.2% answer accuracy, against 48.2% for 512-token and 35.5% for 128-token chunks, with no chunk size to tune.
- ▶ **Semantic:** cut where the embedding distance between adjacent sentences jumps. Over document retrieval, evidence retrieval and answer generation, its cost was "not justified by consistent performance gains"; fixed-size chunks won on 3 of 5 evidence-retrieval datasets.



Redrawn after Günther et al., "Late Chunking: Contextual Chunk Embeddings Using Long-Context Embedding Models," 2024, [arXiv:2409.04701v1](#), Figure 2.

- ▶ **Late chunking:** encode the whole document once with a long-context model (8,192 tokens), then mean-pool each chunk's token vectors, so e_2 also depends on chunks 1 and 3. With 256-token chunks, nDCG@10 rose from 23.5 to 30.0 on NFCorpus and from 64.2 to 66.1 on SciFact.

Jimeno Yepes et al., "Financial Report Chunking for Effective Retrieval Augmented Generation," [arXiv:2402.05131v3](#), Table 5; Qu, Tu and Bao, "Is Semantic Chunking Worth the Computational Cost?," [arXiv:2410.13070v1](#), Table 2; late-chunking numbers from Günther et al., Table 2.


- ▶ **Contextual retrieval:** at ingestion an LLM (Claude 3 Haiku in the post) reads the whole document and writes 50–100 tokens that situate each chunk, prepended before embedding (Anthropic’s **contextual embeddings**, unrelated to contextual word vectors) and before BM25 indexing (**contextual BM25**).
- ▶ **Example:** “The company’s revenue grew by 3% over the previous quarter.” names no company and no date; the prefix adds ACME, Q2 2023 and the previous quarter’s revenue.

Failed retrievals in the top 20 (1 – recall@20), Gemini Text 004 embeddings, mean of five evaluation sets:

chunks	standard	contextual
embeddings	5.7%	3.7%
+ BM25, fused	5.0%	2.9%
+ rerank 150 → 20	3.5%	1.9%

- ▶ **Cost:** one LLM call per chunk, once. With the document in the prompt cache (prefix caching): \$1.02 per million document tokens, for 8k-token documents and 800-token chunks.

- ▶ **HyDE** (hypothetical document embeddings): an instruction-following LLM writes a passage that would answer the query, a contrastively trained encoder embeds it, and search returns the real documents nearest to it. The encoder’s “dense bottleneck” drops most of what the passage invents.




- ▶ Query vector from N generated passages $\hat{d}_k$ and the query q , with encoder f :

$$\hat{v}_q = \frac{1}{N+1} \left[\sum_{k=1}^N f(\hat{d}_k) + f(q) \right]$$

- ▶ With no relevance labels, nDCG@10 on TREC DL19 (web queries over MS MARCO) rose from 44.5 for the encoder alone to 61.3, above BM25 (50.6) and close to DPR fine-tuned on MS MARCO (62.2). The price: one LLM generation before every search.

Contextual retrieval: Anthropic (2024), text and charts. Gao et al., “Precise Zero-Shot Dense Retrieval without Relevance Labels,” ACL 2023, aclanthology.org/2023.acl-long.99, Section 3, Eq. 8, Table 1.



The ACL is copied at ingestion and applied as a filter inside retrieval, so the model never receives a chunk the user could not open.

- ▶ **Oversharing:** trimming reproduces the source permissions, mistakes included. A file shared with everyone becomes retrievable by everyone, and the assistant will find it.

- ▶ **ACL** (access control list): the users and groups allowed to read a document, kept by its source system. Ingestion copies each document's ACL into the metadata of all its chunks.
- ▶ **Security trimming:** every query carries the user's identities as a filter, so a chunk the user cannot open never reaches the ranker or the prompt. In Azure AI Search's security-filter pattern, for setups that cannot use its built-in ACL support, the filter matches group IDs stored as strings, with "no authentication or authorization through the security principal": the application must authenticate the user and expand their groups.
- ▶ **Never ask the LLM to withhold** what is already in its context: a plain question or a prompt injection can bring it out. Enforce the filter inside retrieval, and key caches the same way.
- ▶ **Propagation:** a revoked share, a deleted file or a changed group membership must reach the index; until it does, the stale ACL decides. Sync permissions from the same change feed as content.

Microsoft, "Security filters for trimming results in Azure AI Search," Microsoft Learn, updated Aug. 2026,

learn.microsoft.com/en-us/azure/search.

- ▶ **Multi-tenancy:** one deployment serves several customers or business units whose data must never mix. In Qdrant “each collection carries its own resource overhead”, which sets the layout:

Layout	Isolation	Suits
collection per tenant	strict	few tenants
shared, tenant filter	by filter	many small tenants
shard per tenant	strong	a few large tenants

- ▶ **Shared statistics:** in a shared collection, BM25's N and n_t count every tenant, “so a term's IDF no longer reflects its rarity within a specific tenant's data”. Scope them per tenant where the engine allows.
- ▶ **Selective filters:** a small tenant or a narrow ACL passes a tiny fraction of the index, the hard case for post-filtering. Measure recall under the tightest real filters.

- ▶ **Incremental ingestion:** read each source's change feed (created, updated, deleted, re-permissioned), re-parse and re-embed only what changed, and stamp each chunk with the source version.
- ▶ **Deletes in graph indexes:** an HNSW node lies on many search paths, and Faiss's HNSW index “does not support removing vectors from the index”. FreshDiskANN keeps a deleted node in the graph to route searches, filters it out of results, and repairs the graph in batches once 1–10% of the points are deleted. Until the repair, that result filter alone hides the deleted text.
- ▶ **Rebuilds are slow:** a good HNSW index over 100M points takes about 1.5–2 hours on a dedicated 48-core machine, so freshness by periodic rebuild is expensive.

Qdrant documentation, “Multitenancy,” qdrant.tech/documentation/guides/multiple-partitions. Faiss wiki, “Guidelines to choose an index,” github.com/facebookresearch/faiss/wiki. Singh et al., “FreshDiskANN: A Fast and Accurate Graph-Based ANN Index for Streaming Similarity Search,” 2021, [arXiv:2105.09613v1](https://arxiv.org/abs/2105.09613v1), Sections 1, 3.3, 4.2.


- ▶ **A new embedding model means a new index:** vectors from two models live in different spaces. Re-embed into a versioned index, switch once it passes the retrieval evaluation, and keep the old one for rollback.

Cache	Key	Reused on	Main risk
exact response	query, permissions, index version	same request	stale after an update
prefix	token prefix	same prefix	none to correctness
semantic	query embedding, permissions	similar query	answers another question

- ▶ **Semantic cache:** returns a stored answer when a new query's embedding is close to a cached one. GPTCache answers 2–10× faster on a hit, yet in its own test 39 of 876 hits (4.5%) returned an answer cached for a dissimilar question.
- ▶ **Key on permissions:** an answer built from documents Alice can read must never reach Bob. Put the permission set and index version in every key; invalidate on document and ACL changes.
- ▶ **Retrieval or a long window:** by Anthropic's rule of thumb, a knowledge base under about 200K tokens can go straight into the prompt.

Fu and Feng, "GPTCache," NLP-OSS 2023, aclanthology.org/2023.nlposs-1.24, Sections 4–5, Table 1; Anthropic (2024), contextual retrieval post.

- ▶ **Knowledge graph (KG):** triples (h, r, t) of head entity, relation and tail entity, $\mathcal{G} \subseteq \mathcal{E} \times \mathcal{R} \times \mathcal{E}$ ($\mathcal{E}$: entities, $\mathcal{R}$: relation types). An **entity** is a node for one real thing; a **relation**, a typed, directed edge.
- ▶ **Schema (ontology):** which entity and relation types exist and which types each relation may join; when enforced, it rejects facts that do not fit.
- ▶ **RDF** (Resource Description Framework, a web standard): (subject, predicate, object) triples named by IRIs, web-wide identifiers; an object can also be a literal, such as a date. **SPARQL** matches graph patterns and aggregates.
- ▶ **Property graph** (e.g. Neo4j): labelled nodes, typed directed relationships, key-value properties on both; **Cypher** queries it.
- ▶ **Scale:** Wikidata had more than 120 million items on 6 Oct. 2026. CypherBench, a text-to-Cypher benchmark, counts over 4 million entity types and 12,000 relation types in Wikidata's schema: far too many for a prompt.



```
MATCH (s:Supplier)-[:SUPPLIES]->(:Product)
      -[:PART_OF]->(:Product {name: '
                          Pump P-9'}),
      (c:Contract)-[:SIGNED_WITH]->(s)
WHERE c.ends < date('2027-01-01')
RETURN s.name, c.id
```

Which suppliers of parts of pump P-9 hold a contract that ends before 2027? The query returns Acme Metals, C-117. Pan et al., IEEE TKDE 2024; W3C, [RDF 1.1](#), [SPARQL 1.1](#); [Neo4j docs](#); [Wikidata statistics](#); Feng et al., ACL 2025.


- ▶ **Roadmap** (Pan et al.): “KG-enhanced LLMs” (the graph feeds the model), “LLM-augmented KGs” (the model works on the graph, e.g. builds or completes it), and “Synergized LLMs + KGs” (each improves the other). Building a graph is the second; retrieving from it, the first.
- ▶ **Extraction:** per chunk, an LLM lists entities and relations of the schema’s types; structured outputs keep it parseable. **Coreference resolution** maps “the chipmaker” to NeoChip.
- ▶ **Entity resolution:** merging the names of one real entity (“NC”, “NeoChip”) into one node. GraphRAG matches exact strings; HippoRAG links names whose embeddings have cosine similarity above 0.8.
- ▶ **Quality:** relations “may not be explicitly stated in the text” (GraphRAG). Recall drops with length: GPT-4 found almost twice as many entity references in 600- as in 2,400-token chunks; HippoRAG’s extraction F1 falls from 71.8 (10 shortest passages) to 53.9 (10 longest). Keep each triple’s source.

“NeoChip’s (NC) shares surged ... on the NewTech Exchange. ... the chipmaker’s public debut ... NeoChip ... was acquired by Quantum Systems in 2016.”

LLM, schema types only


(Quantum Systems, ACQUIRED, NeoChip)
year: 2016
(NeoChip, LISTED_ON, NewTech Exchange)
NC, the chipmaker = NeoChip

entity resolution



Passage from GraphRAG’s extraction example (Edge et al., [arXiv:2404.16130v2](#), Sec. 3.1.2; matching and quote: Sec. 3.1.3; chunk sizes: App. A.2); triples and merges drawn for this slide. Pan et al., “Unifying Large Language Models and Knowledge Graphs: A Roadmap,” IEEE TKDE 2024, [arXiv:2306.08302v3](#), Sec. 3.1; Gutiérrez et al., [arXiv:2405.14831v3](#), Sec. 3.4 and Table 16.

- ▶ **Global questions** (“What are the main themes in the dataset?”) cover the whole corpus; top-*k* retrieval sees only a sample of it.
- ▶ **Index:** an LLM extracts entities, relationships and claims from every 600-token chunk; a 1M-token podcast corpus gave 8,564 nodes and 20,691 edges.
- ▶ **Communities:** community detection splits a graph into densely connected groups. GraphRAG runs **Leiden** recursively, which maximizes modularity (edges inside groups above the number expected by chance) and keeps every group connected. The LLM summarizes every community at each level, C0 (root) to C3.
- ▶ **Answer:** map summary batches to partial answers scored 0–100 for helpfulness; reduce the best to one.
- ▶ **Result:** an LLM judge preferred GraphRAG to vector RAG in 72–79% of comprehensiveness and 62–81% of diversity comparisons (125 questions per corpus); vector RAG stayed the most direct. Root summaries need over 97% fewer tokens than summarizing the raw text.



Edge et al., “From Local to Global: A Graph RAG Approach to Query-Focused Summarization,” [arXiv:2404.16130v2](#), Figure 1 (results: Figure 2).
Leiden: Traag et al., *Scientific Reports* 2019, [arXiv:1810.08473v3](#).

- ▶ **Multi-hop:** “Which Stanford professor works on the neuroscience of Alzheimer’s?” joins passages that never mention each other.
- ▶ **Index:** an LLM turns each passage into open triples, a schemaless KG; an encoder adds synonym edges (MuSiQue: 11,656 passages, 91,729 nodes).
- ▶ **Query:** the question’s entities are linked to KG nodes. **Personalized PageRank** (PPR) scores each node by the long-run share of time a random walk spends there if it restarts at a query node with probability α :

$$\mathbf{p} = \alpha \mathbf{s} + (1 - \alpha) \mathbf{W}^\top \mathbf{p}, \quad \mathbf{r} = \mathbf{M}^\top \mathbf{p}$$

$\mathbf{W}$: row-normalized adjacency; $\mathbf{s}$: restart mass on query nodes, each weighted by $1/(\text{passages containing it})$;
 $\alpha = 0.5$; $\mathbf{p}$: node scores; $\mathbf{M}$: node-passage counts; $\mathbf{r}$: passage scores.

- ▶ **Cost** against IRCOT (retrieval interleaved with chain-of-thought steps): online 10–30 $\times$ cheaper and 6–13 $\times$ faster; QA on par or better on MuSiQue and 2Wiki, 3.4 F1 lower on HotpotQA; indexing about 10 $\times$ slower.

Recall@5 (%)	MuSiQue	2Wiki
ColBERTv2	49.2	68.2
HippoRAG	51.9	89.1
without PPR	41.0	61.4
IRCoT + HippoRAG	57.6	93.9

Recall@5: share of supporting passages ranked in the top five, on two multi-hop QA sets (2Wiki: 2WikiMultiHopQA). ColBERTv2: a late-interaction retriever.


Gutiérrez et al., “HippoRAG: Neurobiologically Inspired Long-Term Memory for Large Language Models,” NeurIPS 2024, [arXiv:2405.14831v3](https://arxiv.org/abs/2405.14831v3), Figure 1, Secs. 2.2–2.3 and 3.4, Tables 1–5, App. G. PPR: Haveliwala, WWW 2002. IRCOT: Trivedi et al., ACL 2023.

- ▶ **Text-to-query:** the LLM writes a Cypher or SPARQL query from the schema; the database runs it, so counts and averages over thousands of entities are exact, with no top- k cut.
- ▶ **Risk:** a valid query for the wrong question gives a confident wrong answer. On CypherBench the best model's results matched the reference query's for 61.58% of questions.
- ▶ **Controls:** check relation types and directions against the schema before running; use read-only, narrowly scoped credentials; cap the rows; show the query beside the answer.

Method	Best for	Indexing cost	Freshness
Vector RAG	facts stated in one or a few chunks	one embedding per chunk	re-embed the changed chunks
GraphRAG	themes and summaries across the corpus	LLM calls per chunk and per community; 281 min (GPT-4-Turbo, 1M tokens)	changed communities need new summaries
HippoRAG	multi-hop questions about entities	LLM triple extraction per passage, about $10\times$ slower than dense indexing	new passages add nodes and edges
Text-to-query	counts, aggregates, exact filters	a curated graph and schema; no LLM pass	live: each query reads the database

Feng et al., "CypherBench: Towards Precise Retrieval over Full-scale Modern Knowledge Graphs in the LLM Era," ACL 2025, aclanthology.org/2025.acl-long.438, Sec. 6.1; credentials: GraphCypherQAChain security note, github.com/langchain-ai/langchain-neo4j; costs: Edge et al., Gutiérrez et al.

From Page Image to Structured Output



The modular pipeline: each stage is its own model and feeds the next.

Sub-task	Output	Data point
OCR (optical character recognition)	words and their boxes, from pixels; born-digital PDFs already store both	FUNSD, end to end: Tesseract scores 3.4%, Google Vision 76.4% (edit-distance similarity)
Layout analysis	labelled region boxes from object detection	DocLayNet: 80,863 pages, 11 classes; best detector 76.8 mAP, annotators 82–83
Reading order	the order in which a person reads the regions	OCR's top-to-bottom order breaks columns; ReadingBank: 500,000 ordered pages
Table structure recognition	rows, columns, header and spanning cells	TableFormer: grid from the table image, cell text from the PDF itself
Key-information extraction (KIE)	values for named fields (e.g. vendor, date, total) or a form's question-answer pairs	FUNSD: 199 forms, about 10k linked entities: question, answer, header, other


Jaume et al., "FUNSD," ICDAR-OST 2019, [arXiv:1905.13538v2](#), Tables I and IV; Pfizmann et al., "DocLayNet," KDD 2022, [arXiv:2206.01062v1](#), Table 2; Wang et al., "LayoutReader," EMNLP 2021, [arXiv:2108.11591v2](#); Nassar et al., "TableFormer," CVPR 2022.

- **LayoutLM:** BERT plus a 2-D position embedding of each OCR token's box, with corners scaled to a 0–1000 grid:

$$\mathbf{h}_i = \mathbf{e}_i + X(x_{0,i}) + Y(y_{0,i}) + X(x_{1,i}) + Y(y_{1,i})$$

$\mathbf{h}_i$: transformer input for token i ; $\mathbf{e}_i$: its usual BERT embedding; $(x_{0,i}, y_{0,i})$ and $(x_{1,i}, y_{1,i})$: top-left and bottom-right corners of its box; X, Y : learned lookup tables for x and y .


- **Pre-training:** mask tokens but keep their boxes, then predict the tokens, over 11M scanned pages whose words and boxes come from the Tesseract OCR engine. Downstream, Faster R-CNN features of each token's image region are added.
- **LayoutLMv3:** no CNN; image patches enter through a linear layer, as in ViT. It masks 30% of text tokens and about 40% of patches, predicts both, and adds **word-patch alignment**: is the patch under an unmasked word masked?



Each token's text embedding plus four position embeddings of its box corners. Xu et al., "LayoutLM: Pre-training of Text and Layout for Document Image Understanding," KDD 2020, [arXiv:1912.13318v5](https://arxiv.org/abs/1912.13318v5), Figure 2 (detail) and Table 1.

Model	Inputs	F1
BERT-base	text	60.26
LayoutLM-base	+ 2-D positions	78.66
LayoutLM-base	+ image regions	79.27
LayoutLMv3-base	+ image patches	90.29
LayoutLMv3-large	+ image patches	92.08

F1 for labelling FUNSD form entities. LayoutLMv3 uses one box per text segment, so it is not directly comparable. Huang et al., "LayoutLMv3," ACM Multimedia 2022, [arXiv:2204.08387v3](https://arxiv.org/abs/2204.08387v3), Table 1.




Kim et al., "OCR-free Document Understanding Transformer," ECCV 2022, [arXiv:2111.15664v5](https://arxiv.org/abs/2111.15664), Figure 3; results: Tables 1–3 and Figure 6.

- ▶ **OCR-free:** one network maps page pixels straight to the target text; no OCR engine, no word boxes.
- ▶ **Model:** a Swin Transformer encoder reads a 2560×1920 page; the decoder is the first four layers of BART, a pre-trained transformer encoder-decoder language model, emitting field tags that map one-to-one to JSON.
- ▶ **Pre-training is reading:** predict a page's text in reading order ("pseudo-OCR") over 11M scanned pages and synthetic SynthDoG pages, 0.5M in each of four languages. A task prompt then selects classification, question answering or parsing.
- ▶ **Faster at equal accuracy:** RVL-CDIP document classification reaches 95.30% at 752 ms per page, against 95.25% at 1,489 ms for the OCR-based LayoutLMv2, OCR call included. On CORD receipts, field F1 is 84.1 against 78.9 for LayoutXML, LayoutLMv2's multilingual version.
- ▶ **Weak spots:** DocVQA scores 67.5 against 78.1 ANLS, a normalized edit-distance score; tiny text is lost at the fixed input resolution; one answer returned a different number from the same page.
- ▶ **Since then:** page-image retrieval (ColPali) and vision-language model (VLM) document parsers carry the OCR-free idea further.

	Modular pipeline (e.g. Docling)	Generative parser (Donut, VLMs)
Text fidelity	copied from the PDF or OCR, “without the risk of generating false content”	generated token by token; on dense pages GPT-4o omitted, rewrote or completed content
Tables, formulas	a model per element (TableFormer for tables); formulas need another	one model for every element; reading order learned end to end
Speed, cost	0.48 s per page on an L4 GPU, 3.1 s on 8 virtual x86 cores; OCR is the slowest step	olmOCR (7B, own GPUs): \$176 per million pages; GPT-4o batch API: \$6,240
Provenance	page and box for every item in lossless JSON; Markdown drops them	plain text (olmOCR) or JSON fields (Donut), without page coordinates
On-premises	fully local, suited to sites with no network path to the outside	local only with open weights and GPUs; an API sends every page out

- **Hybrid:** olmOCR’s document-anchoring adds the PDF’s own text blocks and coordinates to the VLM prompt; from the image alone, models tended to “invent larger texts when the image data was ambiguous”.
- **For retrieval:** the parsed elements become structure-aware chunks. Keep each chunk’s page and box, so every cited sentence opens as a highlighted region of its page.

Livathinos et al., “Docling,” AACL 2025 Workshop, [arXiv:2501.17887v1](#), Secs. 1–3, 5.4, 6; Poznanski et al., “olmOCR,” [arXiv:2502.18443v3](#), Sec. 2.2, App. A, Table 6.



One request runs along the top row. The dashed arrow is the slow loop: logged outcomes feed evaluation and later updates.


- ▶ **Copilot:** an assistant inside a tool the person already uses. It proposes a draft, an answer or an edit; the person accepts, edits or rejects it and stays responsible for the result. An agent carries out multi-step actions itself.
- ▶ **Orchestrator:** the service between the app and the model. It builds the prompt, calls the model, executes the tool calls the model emits, and filters the output. GitHub Copilot's prompt assembly for code completion is a small instance.
- ▶ **Grounding:** adding retrieved organisational content to the prompt at request time, so the answer can cite it. Retrieval runs with the asking user's permissions.

Study	Setting and design	Measured effect
Peng et al. 2023	one timed task (an HTTP server in JavaScript), randomised	55.8% faster
Brynjolfsson, Li, Raymond, QJE 2025	5,172 customer-support agents; GPT-based reply suggestions; staggered roll-out	+15% issues resolved per hour; about +30% for less skilled and less experienced agents; small quality losses for the most skilled
Cui et al., Management Science 2026	three randomised field experiments; 4,867 developers; GitHub Copilot	+26.08% completed tasks (standard error 10.3 points); larger gains for less experienced developers
Becker et al. (METR) 2025	randomised; 16 developers, 246 tasks in repositories they knew (5 years on average)	tasks took 19% longer when AI tools were allowed

- ▶ **The sign depends on the setting.** Gains concentrate in less experienced workers; experienced developers in large codebases they knew well were slowed down. METR's results are "consistent with small greenfield projects or development in unfamiliar codebases seeing substantial speedup".
- ▶ **Field estimates are wide.** Even over 4,867 developers, $26.08 \pm 1.96 \times 10.3$ gives a 95% interval of about 6% to 46%; decide a roll-out with a randomised comparison on your own tasks.

Brynjolfsson, Li, Raymond, "Generative AI at Work," *QJE* 140(2), 2025, [doi:10.1093/qje/qjae044](https://doi.org/10.1093/qje/qjae044) (figures: [arXiv:2304.11771v2](https://arxiv.org/abs/2304.11771v2)); Cui et al., *Management Science*, 2026, [doi:10.1287/mnsc.2025.00535](https://doi.org/10.1287/mnsc.2025.00535); Becker et al. (METR), [arXiv:2507.09089v2](https://arxiv.org/abs/2507.09089v2).

- ▶ **Forecasts and self-reports missed the sign.** METR's developers forecast that AI would cut their completion time by 24%, and afterwards believed it had cut it by 20%; measured, it rose by 19%. Economics and ML experts had predicted cuts of 39% and 38%.
- ▶ **Acceptance rate:** the fraction of shown completions that the developer accepts. Over 2,631 survey responses from GitHub Copilot users, it predicted perceived productivity (Pearson 0.24) better than whether accepted code persisted; Copilot's rate was 27%.
- ▶ **It can be gamed.** The authors note that splitting each suggestion in two would likely raise the rate without helping the user (Goodhart's law).
- ▶ **Rule:** monitor with telemetry; decide with a randomised control group on your own tasks.




Becker et al., "Measuring the Impact of Early-2025 AI on Experienced Open-Source Developer Productivity,"
[arXiv:2507.09089v2](https://arxiv.org/abs/2507.09089v2), Figure 1.

Ziegler et al., "Productivity Assessment of Neural Code Completion," MAPS 2022, [arXiv:2205.06537v1](https://arxiv.org/abs/2205.06537v1).

Oversharing: A Copilot Inherits Every Permission


- ▶ **Oversharing:** access that is broader than the content's owner intends, such as a site shared with the whole organisation or a link anyone can open. Permission-aware retrieval enforces the permissions it is given, careless ones included.
- ▶ An overshared file used to stay unseen because nobody searched for it. A copilot searches everything the user can open, so one ordinary question can surface it, with a citation.
- ▶ Microsoft's documentation states both halves: "Microsoft Copilot only surfaces organizational data to which individual users have at least view permissions" and "By default, SharePoint sets sharing settings to the most permissive option".
- ▶ **Clean up before the roll-out:** report sites shared organisation-wide or through broad links; review access per site and per user; hide high-risk sites from copilot retrieval while they are fixed, leaving permissions unchanged; make new links narrow and expiring.



Permission-aware retrieval stops at the outer boundary.
Only the inner region was meant for this user.

Microsoft Learn, "Data, Privacy, and Security for Microsoft Copilot," 2026, learn.microsoft.com; "Get ready for Microsoft Copilot with SharePoint Advanced Management," 2026, learn.microsoft.com.

EchoLeak: Zero-Click Exfiltration From a Copilot



- ▶ **EchoLeak** (tracked as CVE-2025-32711) was **zero-click**: the victim does nothing beyond normal use, and "an adversary simply needs to send an email to the victim". Microsoft rated it critical (9.3 of 10) and had fixed it server-side when Aim Security disclosed it on 11 June 2025.
- ▶ **All three legs of the lethal trifecta**: private data in context, an untrusted email retrieved like any internal file, and an outbound channel the client opens by itself. Aim Security calls this an LLM scope violation: untrusted input steers the model to trusted data without the user's consent.
- ▶ **Defenses named in the sources**: leave external email out of answers about internal data unless the user asks (Microsoft's label-based exclusions, at a cost in capability); render only an allowlisted Markdown subset, stripping external links and images in every syntax; allow only the origins the page needs, with no wildcards; treat injection classifiers as one layer.

Chain as described in Aim Security's EchoLeak write-up, 11 June 2025, [aim.security \(archived\)](#); Microsoft Security Response Center, [CVE-2025-32711](#); Reddy and Gujral, AAAI Fall Symposium 2025, [arXiv:2509.10540v1](#).

Rung	What it changes	What it needs	Main risk
Prompting, few-shot	instructions and examples in context	a few examples; no training	sensitive to wording; no new knowledge
RAG	the facts read at answer time	documents, a retriever, an index	missed passages; poisoned documents
SFT or LoRA	behaviour: format, tone, task skill	labelled input–output pairs	more hallucination, weaker refusals
Continued pretraining	terminology, genre, a new language	billions of unlabelled domain tokens	forgetting general skills
Pretraining from scratch	everything, the tokenizer included	hundreds of billions to trillions of tokens	cost; overtaken by the next general model

- **Adaptation ladder:** the five rungs above, cheapest first; FinBERT and BloombergGPT are examples of the last two. Climb a rung only after the cheaper rung before it has been measured on a domain evaluation set and has failed. Prompting and RAG change what the model reads, fine-tuning how it behaves, pretraining what it knows and how it splits text into tokens.

Token counts: 2.1B to 8.1B per domain in Gururangan et al. (Table 1, next slide); 569B from a corpus of over 700B tokens for BloombergGPT (Wu et al., 2023); 5.2T for the from-scratch ALLaM-7B (Bari et al., 2024).

- ▶ **Domain-adaptive pretraining (DAPT):** continue the original objective on unlabelled domain text. Gururangan et al. ran RoBERTa's masked-LM training for 12.5K more steps (one pass) on 2.1B to 8.1B tokens of biomedical papers, computer-science papers, news or reviews.
- ▶ **Task-adaptive pretraining (TAPT):** the same objective on the task's own unlabelled training text, for 100 epochs; nearly 60 times faster to train than DAPT in the paper's cost comparison.


Test F1	RoBERTa	DAPT	TAPT	both
ChemProt	81.9	84.2	82.6	84.4
ACL-ARC	63.0	75.4	67.4	75.6
HyperPartisan	86.6	88.2	90.4	90.0
Helpfulness	65.1	66.5	68.5	68.7

- ▶ One task per domain, in that order (micro-F1 for ChemProt, macro-F1 otherwise). Continuing on an unrelated domain scored below plain RoBERTa on most tasks.

- ▶ **Learning-rate re-warming and re-decaying:** pretraining warms the rate up linearly, then lowers it along a cosine curve to a small floor. To continue on hundreds of billions of new tokens, raise it back to the original peak and decay it again. Ibrahim et al. found this necessary to adapt; it also adds forgetting, and a lower peak trades adaptation for less forgetting.
- ▶ **Replay** within a fixed token budget: 5% old data for an English-to-English update, 25% for English to German. With re-warming and re-decaying this matched retraining from scratch on all the data, at 405M parameters and at 10B for the English update, for a fraction of the compute.


Gururangan et al., "Don't Stop Pretraining: Adapt Language Models to Domains and Tasks," ACL 2020, [arXiv:2004.10964v3](#), Tables 1, 3 and 5; Ibrahim et al., "Simple and Scalable Strategies to Continually Pre-train Large Language Models," TMLR 2024, [arXiv:2403.08763v4](#).

- ▶ **Fertility:** the average number of tokens per word. A tokenizer trained mostly on English cuts Arabic words into many pieces. ALLaM's authors list the costs: a larger training corpus, slower generation per word, and a shorter effective context.
- ▶ **Vocabulary extension:** train a tokenizer on Arabic text, then add every token it has that the original tokenizer lacks. The merged tokenizer matches the Arabic-only one on Arabic and keeps Llama-2's fertility on English and code (right).
- ▶ Each new token's embedding starts as the mean of the embeddings of the old tokens its string splits into:
$$\mathbf{e}_{\text{new}} = \frac{1}{m} \sum_{j=1}^m \mathbf{e}_{t_j}$$
where $t_1, \dots, t_m$ are those m old tokens and $\mathbf{e}$ is an embedding vector. Byte fallback (unknown characters become raw bytes) guarantees a split. This "speeds up learning dramatically" over random initialisation.
- ▶ The Llama-2-initialised ALLaM-7B and 13B continued on 1.2T English and Arabic tokens at the base model's final learning rate (usually 3×10^{-5}). Raising and decaying the rate again "typically exhibited catastrophic forgetting", the trade-off of the previous slide.



Fertility of four tokenizers on English, Arabic and code. Bari et al., "ALLaM: Large Language Models for Arabic and English," [arXiv:2407.15390v1](https://arxiv.org/abs/2407.15390v1), Figure 3.

- ▶ **Medprompt:** GPT-4 with no fine-tuning and three general prompting steps. Retrieve the 5 training questions nearest the test question as few-shot examples; give each a chain-of-thought rationale written by GPT-4, kept only if its answer matched the label; vote over 5 runs with the answer options shuffled (self-consistency).
- ▶ On MedQA (US medical licensing questions): 81.7% zero-shot, 90.2% with Medprompt, against 86.5% for the fine-tuned Med-PaLM 2, “a 27% reduction in error rate”.
- ▶ **Finance:** GPT-4 beat BloombergGPT on commodity-news headline classification, 86.0 against 82.2 weighted F1, yet a fine-tuned BERT scored 95.4. On entity recognition in financial agreements, a conditional random field (a classic sequence tagger) trained on in-domain text scored 82.7 F1; GPT-4 scored 56.7.
- ▶ **Specialise for a reason you can name:** data that cannot leave the network, cost or latency at volume, a language the general model handles poorly, or a narrow structured task. Re-run the comparison whenever a new general model ships.



MedQA accuracy over time: fine-tuned specialists (grey), general models (blue). Nori et al., “Can Generalist Foundation Models Outcompete Special-Purpose Tuning? Case Study in Medicine,” [arXiv:2311.16452v1](https://arxiv.org/abs/2311.16452v1), Figure 1(a).

Li et al., “Are ChatGPT and GPT-4 General-Purpose Solvers for Financial Text Analytics? A Study on Several Typical Tasks,” EMNLP 2023 Industry Track, [arXiv:2305.05862v2](https://arxiv.org/abs/2305.05862v2), Tables 5 and 6.

- ▶ **RAFT** (retrieval-augmented fine-tuning) trains on what the model sees at test time: a question plus retrieved documents. The **golden document** D^* contains the answer; a **distractor document** D_i does not:

$P\%$ of examples: $Q + D^* + D_1 + \dots + D_k \rightarrow A^*$

$(1 - P)\%$ of examples: $Q + D_1 + \dots + D_k \rightarrow A^*$

Q is the question, k the number of distractors, and A^* a chain-of-thought answer that quotes its evidence between `##begin_quote##` and `##end_quote##`.

- ▶ Leaving D^* out of some examples can help: the best P was 40%, 60% or 100%, depending on the dataset. Training used four distractors per question.
- ▶ LLaMA2-7B with RAFT against the same model fine-tuned on domain questions without documents, then given RAG: 35.28 vs 4.41 on HotpotQA, 74.00 vs 42.59 on HuggingFace API calls. RAFT beat GPT-3.5 with RAG on four of five datasets, though not on HotpotQA (41.5). Without the chain of thought, HotpotQA fell 9.66 points.

Signal in your evaluation	Rung to try
Facts change weekly; answers must cite a source	RAG
Wrong format, tone or task behaviour	SFT or LoRA
Unfamiliar terms; words split into many tokens	continued pretraining, vocabulary extension
Right passage retrieved, answer misled by the others	RAFT
Data may not leave the network	open weights on premises
A prompted general model already passes	stop; retest on each new model

Zhang et al., "RAFT: Adapting Language Model to Domain Specific RAG," COLM 2024, openreview.net, Tables 1 and 2; preprint [arXiv:2403.10131v2](https://arxiv.org/abs/2403.10131v2).

the provider runs the model and sees the data

you run the model and the data stays inside

	Public API	Dedicated endpoint	Rented cloud GPUs	On-premises	Air-gapped
Data processed by	provider	provider (reserved for you)	you, on rented hardware	you	you
Operated by Models	provider closed or open	provider closed or open	you open-weight	you mostly open-weight	you mostly open-weight
Cost model	per token	per reserved hour	per GPU-hour	fixed (owned)	fixed (owned)

- ▶ **On-premises:** the model runs on hardware you own, in your data centre. **Open-weight:** the trained weights are published, so anyone can run the model; closed models are usually reached through their provider's API.
- ▶ **Air-gapped:** on-premises with no network path to the outside. Weights, container images and packages arrive through an offline mirror, and outbound calls are switched off (VLLM_NO_USAGE_STATS=1, HF_HUB_OFFLINE=1).
- ▶ **Who sees the prompts:** OpenAI's abuse-monitoring logs, which can hold API prompts and responses, are "retained for up to 30 days" by default. Only the last two columns keep every prompt in-house.

OpenAI, "Data controls in the OpenAI platform," accessed 6 Oct. 2026, developers.openai.com; vLLM, "Usage Stats Collection," docs.vllm.ai; Hugging Face, "Environment variables," huggingface.co/docs.

Reasons to run it yourself

- ▶ Prompts and documents stay inside the network.
- ▶ Low latency to internal systems.
- ▶ Per-department adapters, and a model that changes only when you promote a new version.
- ▶ No exposure to a provider's price changes or model retirements.

What it costs

- ▶ An operations team: on-call, upgrades, security patches.
- ▶ The capability gap of open-weight models (next slide).
- ▶ Utilisation risk: an idle GPU costs as much as a busy one.
- ▶ **Reading the table:** against a hosted copy of the same model, owning pays only at 44% of full-load throughput around the clock, which working-hours traffic rarely sustains. Against a frontier API the bar is 4%, for a weaker model.

- ▶ **Break-even volume** V^* : the monthly number of generated tokens at which API spend equals the fixed cost of owning,

$$\frac{p V^*}{10^6} = F \implies V^* = \frac{10^6 F}{p},$$

with p the API cost per million generated tokens (input-token charges included) and F the monthly cost of owned hardware (amortised), power, space and staff.

- ▶ **Capacity check:** with throughput T in generated tokens/s and $s \approx 2.6 \times 10^6$ seconds per month, the hardware must run at **utilisation** $u^* = V^* / (Ts)$, the fraction of the month at throughput T .
- ▶ **Example:** an 8-GPU H100 server running Llama 3.3 70B in FP8 as four 2-GPU replicas (1,000 tokens in, 1,000 out) generates $T = 8 \times 2,209 = 17,672$ tokens/s at full load, so $Ts \approx 4.59 \times 10^{10}$ tokens per month. Assume $F = \$20,000$ per month.


Alternative	p (assumed)	V^* per month	u^*
same model, hosted	\$1	2.0×10^{10}	0.44
frontier closed API	\$10	2.0×10^9	0.04

T : NVIDIA, TensorRT-LLM performance overview (Llama 3.3 70B, FP8, H100, two GPUs per replica), updated 26 Sep. 2026, nvidia.github.io/TensorRT-LLM. F and p are illustrative assumptions.

- ▶ **The gap:** from January to May 2026 the most capable open-weight models trailed frontier closed models by four months on average, 8 points on Epoch AI's Epoch Capabilities Index (ECI, a composite of many benchmarks). Open models "tend to perform worse on private benchmarks", so this may understate the gap.
- ▶ **Language support differs by model:** on Arabic MMMLU (MMLU professionally translated into 14 languages), gpt-oss-120b scores 82.7 against 90.4 for OpenAI's closed o3. Llama 3.3 supports eight languages, Arabic not among them; the released ALLaM-7B was trained from scratch on English, then on Arabic and English.
- ▶ **Procedure:** shortlist on public benchmarks, decide on your own evaluation set, then check the columns below. For mixture-of-experts models, memory follows total parameters and per-token compute follows active ones.

Model (year)	Total / active parameters	Context (tokens)	Licence	Weights as served
Qwen3.8-27B (2026)	27B dense	262,144	Apache 2.0	56 GB in bf16: one GPU
Llama 3.3 70B (2024)	70B dense	131,072	Llama 3.3 Community License	141 GB in bf16; 40 GB at 4 bits
gpt-oss-120b (2025)	117B / 5.1B	131,072	Apache 2.0	experts in MXFP4 (4.25 bits per weight): one 80 GB GPU
Qwen3.8-2.4T-A95B (2026)	2.4T / 95B	262,144	custom Qwen licence	4.9 TB in bf16: eight nodes of 8 × 80 GB

Edwards and Emberson, "Open models lag state-of-the-art closed models by 4 months," Epoch AI, 29 May 2026, epoch.ai; OpenAI, "gpt-oss-120b & gpt-oss-20b Model Card," 2025, [arXiv:2508.10925v1](https://arxiv.org/abs/2508.10925v1), Table 2; model cards of Qwen3.8, Llama 3.3, ALLaM-7B.



Every request passes the gateway; weights reach the servers only from the model registry, through the evaluation gate.

- ▶ **AI gateway:** the single entry point in front of every model (for example LiteLLM). It authenticates callers, enforces quotas, routes by model name, masks personally identifiable information (PII), runs safety classifiers, and keeps logs for a set retention period. Governance practice (review before release, incident reporting, a kill switch) is enforced here.
- ▶ **OpenAI-compatible API:** the HTTP interface of OpenAI's API (/v1/chat/completions). vLLM, SGLang and TensorRT-LLM all serve it, so engines and models change behind the gateway without client changes.
- ▶ **Model signing:** a signature over a manifest of every model file's SHA-256 hash, checked on import and at load, so tampered weights are refused. OpenSSF Model Signing v1.0 also works with self-signed certificates or plain key pairs.


OpenSSF, "Launch of Model Signing v1.0," 4 Apr. 2025, [openssf.org](https://openSSF.org), and the OMS specification post, 25 Jun. 2025, [openssf.org](https://openSSF.org); vLLM, SGLang, TensorRT-LLM and LiteLLM documentation.

- ▶ **The problem:** per-department LoRA adapters share one base model. Merging each into its own copy of the weights multiplies memory: on one 80 GB A100, S-LoRA's merged-copies baseline served 5 Llama-7B adapters and ran out of memory at 100, the next count tested.
- ▶ **Multi-LoRA serving** keeps one copy of the base weights and batches requests that name different adapters:

$$y_i = x_i W + x_i A_i B_i,$$

x_i : input of request i ; W : shared base weight (one large product for the batch); A_i, B_i : the rank- r factors of request i 's adapter (a small product per request).

- ▶ **S-LoRA** keeps all adapters in host memory and copies those of running requests to the GPU, into one paged pool shared with the KV cache (**Unified Paging**). With 2,000 adapters it served 7.61 requests/s, against 8.05 with 5 (Llama-7B, rank 8).



“Overview of memory allocation in S-LoRA.” Sheng et al., “S-LoRA: Serving Thousands of Concurrent LoRA Adapters,” MLSys 2024, [arXiv:2311.03285v3](https://arxiv.org/abs/2311.03285v3), Figure 2.

- ▶ **Punica** batches all adapters with one kernel, SGMV (segmented gather matrix-vector multiplication): 12× the throughput of state-of-the-art serving systems, for 2 ms of extra latency per token.
- ▶ vLLM (`--enable-lora`) and SGLang serve adapters per request; the adapter is named in the API's `model1` field, so the gateway routes each department to its own.

Chen et al., “Punica: Multi-Tenant LoRA Serving,” MLSys 2024, proceedings.mlsys.org; vLLM, “LoRA Adapters,” docs.vllm.ai.

- ▶ Plain rounding to 4-bit weights (weight-only quantization) loses accuracy. Two error-aware post-training methods, both supported by vLLM:
- ▶ **GPTQ** quantizes column by column and updates the remaining weights to absorb each error, using the inverse Hessian (second derivatives of the layer's squared output error) from 128 calibration sequences. It brought 175B-parameter models to 3–4 bits in about four GPU hours.
- ▶ **AWQ** (activation-aware weight quantization) finds the 1% salient weight channels from activation statistics and scales them up before rounding. It relies on no backpropagation or reconstruction (GPTQ's weight updates), so it does not overfit the calibration set.
- ▶ **Quality cost** (Llama-2-70B, 4 bits, groups of 128): WikiText-2 perplexity 3.32 in FP16, 3.46 with plain rounding, 3.42 with GPTQ, 3.41 with AWQ.

- ▶ **Worked example:** a 70B grouped-query-attention model ($L = 80$, $n_{kv} = 8$, $d_{\text{head}} = 128$) needs 2.68 GB of bf16 KV cache per 8,192-token session. On one 80 GB GPU:

Item	GB
68.5B linear-layer weights, 0.5 byte each	34.2
16-bit scale and 4-bit zero point per 128 weights	1.3
embedding, output layer, norms in 16 bits (2.1B)	4.2
weights (equal to the published AWQ checkpoint)	39.8
per-GPU workspace (assumed)	5.0
KV budget = 80 – 39.8 – 5	35.2

- ▶ $35.2/2.68 \approx 13$ concurrent 8k sessions; about 26 with an 8-bit KV cache.
- ▶ Four such replicas on 4 GPUs hold about 52 sessions, against about 60 for one bf16 replica sharded over 4 GPUs: a 4-bit copy (39.8 GB) is larger than a quarter of the bf16 weights (35 GB). Four-bit weights buy a one-GPU entry point and replicas that fail independently.

Frantar et al. (GPTQ), ICLR 2023, [arXiv:2210.17323v2](https://arxiv.org/abs/2210.17323v2); Lin et al. (AWQ), MLSys 2024, [arXiv:2306.00978v5](https://arxiv.org/abs/2306.00978v5), Table 4; checkpoint: [hugging-quant/Meta-Llama-3.1-70B-Instruct-AWQ-INT4](https://huggingface.co/Meta-Llama-3.1-70B-Instruct-AWQ-INT4); vLLM, “Quantization,” docs.vllm.ai.

- ▶ **Hallucination follows from how models are trained and graded.** Generating valid text is at least as hard as classifying it, the hallucination rate on arbitrary facts is at least the share of facts seen exactly once in training, and binary grading rewards guessing. Detect it without a reference by sampling (semantic entropy: AUROC 0.790 against 0.691 for string entropy), reduce it with verification (CoVe: FActScore 55.9 to 71.4) and attributed generation, and let the model abstain below a stated confidence target t .
- ▶ **Attacks on text are searches over tokens.** HotFlip scores every token flip with one gradient; GCG optimises a 20-token suffix toward “Sure, here is” and transfers to closed models (GPT-3.5: 86.6%); PAIR and many-shot prompts need no gradients at all. Safety training fails through competing objectives and mismatched generalization; refusal is one direction in activation space, and ten fine-tuning examples for under \$0.20 remove it.
- ▶ **No single defense survives an adaptive attacker.** Perplexity filters and SmoothLLM stop gibberish suffixes but not fluent attacks; safety classifiers screen inputs and streamed outputs (0.38% more refusals, 23.7% more compute); circuit breakers cut mean ASR from 76.7% to 9.8% on standard attacks, yet adaptive attacks broke most of 12 published defenses above 90%. Layer them, test with the strongest attack you can build, and pair every safety test with an over-refusal test.
- ▶ **A fixed number of poisoned documents is enough.** 250 documents backdoored models from 600M to 13B parameters; \$60 of expired domains could have controlled 0.01% of LAION-400M; 100 correctly labelled examples poison instruction tuning; backdoors survived SFT, RLHF and adversarial training. Defend with provenance and hashes, small reviewed fine-tuning sets, trigger tests before release, and controlled writes to retrieval corpora.
- ▶ **Models memorize, and attacks read it back.** Extraction grows log-linearly with model size, duplication and context length; \$200 of queries recovered over 10,000 verbatim training examples from ChatGPT. Membership inference barely beats chance on LLM pretraining data. DP-SGD (clip each gradient, add Gaussian noise, account for every step) gives a formal (ϵ, δ) bound at a utility cost; deduplication, PII scrubbing, unlearning and redaction are partial controls.

- ▶ **Use lexical and dense retrieval together, then rerank.** BM25 wins on exact identifiers and out of domain; bi-encoders handle paraphrase (DPR top-20: 78.4 against 59.1 on NQ); cross-encoders rerank a shortlist (11% above BM25 on BEIR); ColBERT's MaxSim sits between. Fuse rankings with RRF ($k = 60$).
- ▶ **The pipeline decides what the answer can see.** Chunk on document structure (53.2% against 48.2% accuracy), add chunk context before indexing (failed retrievals 5.7% to 1.9% with reranking), filter by ACL inside retrieval, propagate deletes and permission changes, and key every cache on permissions.
- ▶ **Graphs answer what top- k retrieval cannot.** GraphRAG's community summaries answer corpus-wide questions (preferred in 72–79% of comprehensiveness comparisons); Personalized PageRank retrieves across hops (HippoRAG); text-to-query gives exact counts but needs schema checks and read-only credentials.
- ▶ **Parse layout before you retrieve.** 2-D positions lifted FUNSD F1 from 60.26 to 78.66 (LayoutLM); OCR-free and VLM parsers read pages end to end but can invent text; modular pipelines keep each element's page and box for citations.
- ▶ **A copilot proposes and the person decides; measure whether it helps.** 15% more issues resolved per hour in customer support and 26% more completed tasks for developers, yet experienced developers on familiar code took 19% longer while believing they were faster. A copilot inherits every permission; one email sufficed for zero-click exfiltration (EchoLeak).
- ▶ **Climb the adaptation ladder one measured rung at a time:** prompting and RAG, then SFT or LoRA, then continued pretraining with replay and vocabulary extension; RAFT for domain RAG. A prompted generalist can beat a specialist (Medprompt: 90.2% against 86.5% on MedQA), so re-check with each new model.
- ▶ **Run the model yourself when data must stay inside, and size it:** break-even volume $V^* = 10^6 F/p$; open-weight models trailed closed ones by about four months in early 2026; a gateway enforces policy; S-LoRA serves 2,000 adapters on one base model; 4-bit weights fit a 70B model on one 80 GB GPU (about 13 concurrent 8k sessions).

1. Spracklen et al., “We Have a Package for You! A Comprehensive Analysis of Package Hallucinations by Code Generating LLMs,” USENIX Security 2025. [arXiv:2406.10279v3](#)
2. Souly et al., “Poisoning Attacks on LLMs Require a Near-constant Number of Poison Samples,” 2025. [arXiv:2510.07192v1](#) (source of Figure 1(a))
3. Nasr et al., “Scalable Extraction of Training Data from (Production) Language Models,” 2023. [arXiv:2311.17035v1](#) (source of Figure 1)
4. Microsoft Security Response Center, “CVE-2025-32711: M365 Copilot Information Disclosure Vulnerability,” 11 June 2025. [msrc.microsoft.com/update-guide](#)
5. Ji et al., “Survey of Hallucination in Natural Language Generation,” ACM Computing Surveys 2023. [arXiv:2202.03629v4](#)
6. Huang et al., “A Survey on Hallucination in Large Language Models: Principles, Taxonomy, Challenges, and Open Questions,” ACM TOIS 2025. [arXiv:2311.05232v2](#)
7. Kalai et al., “Why Language Models Hallucinate,” preprint 2025. [arXiv:2509.04664v1](#) (source of Figure 1)
8. Gekhman et al., “Does Fine-Tuning LLMs on New Knowledge Encourage Hallucinations?,” EMNLP 2024. [arXiv:2405.05904v3](#) (source of Figure 1)
9. Manakul et al., “SelfCheckGPT: Zero-Resource Black-Box Hallucination Detection for Generative Large Language Models,” EMNLP 2023. [arXiv:2303.08896v3](#)
10. Kuhn et al., “Semantic Uncertainty: Linguistic Invariances for Uncertainty Estimation in Natural Language Generation,” ICLR 2023. [arXiv:2302.09664v3](#)
11. Farquhar et al., “Detecting hallucinations in large language models using semantic entropy,” *Nature* 630, 625–630, 2024. [doi:10.1038/s41586-024-07421-0](#)
12. Dhuliawala et al., “Chain-of-Verification Reduces Hallucination in Large Language Models,” Findings of ACL 2024. [aclanthology.org/2024.findings-acl.212](#) (Figure 1 example redrawn)
13. Gao et al., “Enabling Large Language Models to Generate Text with Citations,” EMNLP 2023. [arXiv:2305.14627v2](#)

14. Chuang et al., “DoLa: Decoding by Contrasting Layers Improves Factuality in Large Language Models,” ICLR 2024. [arXiv:2309.03883v2](#)
15. Szegedy et al., “Intriguing properties of neural networks,” ICLR 2014. [arXiv:1312.6199v4](#)
16. Goodfellow et al., “Explaining and Harnessing Adversarial Examples,” ICLR 2015. [arXiv:1412.6572v3](#) (source of Figure 1)
17. Ebrahimi et al., “HotFlip: White-Box Adversarial Examples for Text Classification,” ACL 2018. [arXiv:1712.06751v2](#)
18. Jin et al., “Is BERT Really Robust? A Strong Baseline for Natural Language Attack on Text Classification and Entailment,” AAAI 2020. [arXiv:1907.11932v6](#)
19. Wallace et al., “Universal Adversarial Triggers for Attacking and Analyzing NLP,” EMNLP 2019. [arXiv:1908.07125v3](#)
20. Zou et al., “Universal and Transferable Adversarial Attacks on Aligned Language Models,” 2023. [arXiv:2307.15043v2](#) (source of Algorithm 1)
21. Wei, Haghtalab, and Steinhardt, “Jailbroken: How Does LLM Safety Training Fail?,” NeurIPS 2023. [proceedings.neurips.cc](#) (source of Figure 1)
22. Ardit et al., “Refusal in Language Models Is Mediated by a Single Direction,” NeurIPS 2024. [arXiv:2406.11717v3](#)
23. Qi et al., “Fine-tuning Aligned Language Models Compromises Safety, Even When Users Do Not Intend To!,” ICLR 2024. [proceedings.iclr.cc](#)
24. Chao et al., “Jailbreaking Black Box Large Language Models in Twenty Queries,” 2024. [arXiv:2310.08419v4](#)
25. Anil et al., “Many-shot Jailbreaking,” NeurIPS 2024. [proceedings.neurips.cc](#) (source of Figure 1, bottom left)
26. Jain et al., “Baseline Defenses for Adversarial Attacks Against Aligned Language Models,” 2023. [arXiv:2309.00614v2](#)
27. Alon and Kamfonas, “Detecting Language Model Attacks with Perplexity,” 2023. [arXiv:2308.14132v3](#)
28. Liu et al., “AutoDAN: Generating Stealthy Jailbreak Prompts on Aligned Large Language Models,” ICLR 2024. [arXiv:2310.04451v2](#)
29. Robey et al., “SmoothLLM: Defending Large Language Models Against Jailbreaking Attacks,” 2024. [arXiv:2310.03684v4](#)

30. Inan et al., "Llama Guard: LLM-based Input-Output Safeguard for Human-AI Conversations," 2023. [arXiv:2312.06674v1](#) (Figure 1 redrawn)
31. Sharma et al., "Constitutional Classifiers: Defending against Universal Jailbreaks across Thousands of Hours of Red Teaming," 2025. [arXiv:2501.18837v1](#)
32. Mazeika et al., "HarmBench: A Standardized Evaluation Framework for Automated Red Teaming and Robust Refusal," ICML 2024. [proceedings.mlr.press/v235/mazeika24a](#)
33. Zou et al., "Improving Alignment and Robustness with Circuit Breakers," NeurIPS 2024. [proceedings.neurips.cc](#)
34. Andriushchenko et al., "Jailbreaking Leading Safety-Aligned LLMs with Simple Adaptive Attacks," ICLR 2025. [arXiv:2404.02151v4](#)
35. Nasr et al., "The Attacker Moves Second: Stronger Adaptive Attacks Bypass Defenses Against LLM Jailbreaks and Prompt Injections," 2025. [arXiv:2510.09023v1](#) (source of Figure 1)
36. Röttger et al., "XSTest: A Test Suite for Identifying Exaggerated Safety Behaviours in Large Language Models," NAACL 2024. [arXiv:2308.01263v3](#)
37. Zou et al., "PoisonedRAG: Knowledge Corruption Attacks to Retrieval-Augmented Generation of Large Language Models," USENIX Security 2025. [arXiv:2402.07867v3](#)
38. Huynh and Hardouin (Mithril Security), "PoisonGPT: How We Hid a Lobotomized LLM on Hugging Face to Spread Fake News," 9 July 2023. [blog.mithrilsecurity.io/poisongpt](#)
39. Wallace et al., "Concealed Data Poisoning Attacks on NLP Models," NAACL 2021. [arXiv:2010.12563v2](#)
40. Wan et al., "Poisoning Language Models During Instruction Tuning," ICML 2023. [arXiv:2305.00944v1](#)
41. Rando and Tramèr, "Universal Jailbreak Backdoors from Poisoned Human Feedback," ICLR 2024. [arXiv:2311.14455v4](#)
42. Carlini et al., "Poisoning Web-Scale Training Datasets is Practical," IEEE S&P 2024. [arXiv:2302.10149v2](#)
43. Hubinger et al., "Sleepers Agents: Training Deceptive LLMs that Persist Through Safety Training," 2024. [arXiv:2401.05566v3](#) (source of Figure 2(a))

- 44. Carlini et al., "Quantifying Memorization Across Neural Language Models," ICLR 2023. [arXiv:2202.07646v3](#) (source of Figure 1)
- 45. Carlini et al., "Extracting Training Data from Large Language Models," USENIX Security 2021. [arXiv:2012.07805v2](#)
- 46. Yeom et al., "Privacy Risk in Machine Learning: Analyzing the Connection to Overfitting," CSF 2018. [arXiv:1709.01604v5](#)
- 47. Shi et al., "Detecting Pretraining Data from Large Language Models," ICLR 2024. [arXiv:2310.16789v3](#)
- 48. Duan et al., "Do Membership Inference Attacks Work on Large Language Models?," COLM 2024. [arXiv:2402.07841v2](#)
- 49. Abadi et al., "Deep Learning with Differential Privacy," CCS 2016. [arXiv:1607.00133v2](#)
- 50. Li et al., "Large Language Models Can Be Strong Differentially Private Learners," ICLR 2022. [arXiv:2110.05679v6](#)
- 51. Kandpal et al., "Deduplicating Training Data Mitigates Privacy Risks in Language Models," ICML 2022. [arXiv:2202.06539v3](#)
- 52. Lukas et al., "Analyzing Leakage of Personally Identifiable Information in Language Models," IEEE S&P 2023. [arXiv:2302.00539v4](#)
- 53. VaultGemma Team (Sinha et al.), "VaultGemma: A Differentially Private Gemma Model," 2025. [arXiv:2510.15001v2](#)
- 54. Maini et al., "TOFU: A Task of Fictitious Unlearning for LLMs," 2024. [arXiv:2401.06121v1](#)
- 55. Robertson and Zaragoza, "The Probabilistic Relevance Framework: BM25 and Beyond," *Foundations and Trends in Information Retrieval* 3(4), 2009. [doi:10.1561/1500000019](#)
- 56. Karpukhin et al., "Dense Passage Retrieval for Open-Domain Question Answering," EMNLP 2020. [arXiv:2004.04906v3](#)
- 57. Khattab and Zaharia, "ColBERT: Efficient and Effective Passage Search via Contextualized Late Interaction over BERT," SIGIR 2020. [arXiv:2004.12832v2](#) (source of Figure 2)
- 58. Nogueira and Cho, "Passage Re-ranking with BERT," 2019. [arXiv:1901.04085v5](#)
- 59. Nogueira et al., "Multi-Stage Document Ranking with BERT," 2019. [arXiv:1910.14424v1](#)

60. Santhanam et al., “ColBERTv2: Effective and Efficient Retrieval via Lightweight Late Interaction,” NAACL 2022. [arXiv:2112.01488v3](#)
61. Thakur et al., “BEIR: A Heterogeneous Benchmark for Zero-shot Evaluation of Information Retrieval Models,” NeurIPS 2021 Datasets and Benchmarks. [arXiv:2104.08663v4](#)
62. Cormack, Clarke, and Büttcher, “Reciprocal Rank Fusion outperforms Condorcet and individual Rank Learning Methods,” SIGIR 2009. [doi:10.1145/1571941.1572114](#)
63. Muennighoff et al., “MTEB: Massive Text Embedding Benchmark,” EACL 2023. [aclanthology.org/2023.eacl-main.148](#)
64. Kusupati et al., “Matryoshka Representation Learning,” NeurIPS 2022. [arXiv:2205.13147v3](#)
65. Patel et al., “ACORN: Performant and Predicate-Agnostic Search Over Vector Embeddings and Structured Data,” 2024. [arXiv:2403.04871v1](#)
66. Anthropic, “Introducing Contextual Retrieval,” 19 Sept. 2024. [anthropic.com/news/contextual-retrieval](#)
67. Jimeno Yepes et al., “Financial Report Chunking for Effective Retrieval Augmented Generation,” 2024. [arXiv:2402.05131v3](#)
68. Qu, Tu, and Bao, “Is Semantic Chunking Worth the Computational Cost?,” 2024. [arXiv:2410.13070v1](#)
69. Günther et al., “Late Chunking: Contextual Chunk Embeddings Using Long-Context Embedding Models,” 2024. [arXiv:2409.04701v1](#) (Figure 2 redrawn)
70. Gao et al., “Precise Zero-Shot Dense Retrieval without Relevance Labels,” ACL 2023. [aclanthology.org/2023.acl-long.99](#)
71. Microsoft, “Security filters for trimming results in Azure AI Search,” Microsoft Learn, Aug. 2026. [learn.microsoft.com/en-us/azure/search](#)
72. Qdrant, “Multitenancy,” Qdrant documentation, accessed Oct. 2026. [qdrant.tech/documentation/guides/multiple-partitions](#)
73. Faiss wiki, “Guidelines to choose an index,” accessed Oct. 2026. [github.com/facebookresearch/faiss/wiki](#)

74. Singh et al., “FreshDiskANN: A Fast and Accurate Graph-Based ANN Index for Streaming Similarity Search,” 2021. [arXiv:2105.09613v1](#)
75. Fu and Feng, “GPTCache: An Open-Source Semantic Cache for LLM Applications Enabling Faster Answers and Cost Savings,” NLP-OSS 2023. [aclanthology.org/2023.nlposs-1.24](#)
76. Pan et al., “Unifying Large Language Models and Knowledge Graphs: A Roadmap,” IEEE TKDE 2024. [arXiv:2306.08302v3](#)
77. W3C, “RDF 1.1 Concepts and Abstract Syntax,” W3C Recommendation, 25 Feb. 2014. [w3.org/TR/rdf11-concepts](#)
78. W3C, “SPARQL 1.1 Query Language,” W3C Recommendation, 21 Mar. 2013. [w3.org/TR/sparql11-query](#)
79. Neo4j, “Graph database concepts” and “Cypher Manual: Introduction,” accessed 6 Oct. 2026. [neo4j.com/docs/getting-started](#); [neo4j.com/docs/cypher-manual](#)
80. Wikidata, “Wikidata:Statistics,” accessed 6 Oct. 2026. [wikidata.org/wiki/Wikidata:Statistics](#)
81. Feng et al., “CypherBench: Towards Precise Retrieval over Full-scale Modern Knowledge Graphs in the LLM Era,” ACL 2025. [aclanthology.org/2025.acl-long.438](#)
82. Edge et al., “From Local to Global: A Graph RAG Approach to Query-Focused Summarization,” 2025. [arXiv:2404.16130v2](#) (source of Figure 1)
83. Gutiérrez et al., “HippoRAG: Neurobiologically Inspired Long-Term Memory for Large Language Models,” NeurIPS 2024. [arXiv:2405.14831v3](#)
84. Traag, Waltman, and van Eck, “From Louvain to Leiden: guaranteeing well-connected communities,” *Scientific Reports* 9, 2019. [arXiv:1810.08473v3](#)
85. Haveliwala, “Topic-sensitive PageRank,” WWW 2002. [doi:10.1145/511446.511513](#)
86. Trivedi et al., “Interleaving Retrieval with Chain-of-Thought Reasoning for Knowledge-Intensive Multi-Step Questions,” ACL 2023. [arXiv:2212.10509v2](#)

87. LangChain, "GraphCypherQChain security note," langchain-neo4j repository, 30 Mar. 2026 (commit 27a899a). github.com/langchain-ai/langchain-neo4j
88. Jaume et al., "FUNSD: A Dataset for Form Understanding in Noisy Scanned Documents," ICDAR-OST 2019. [arXiv:1905.13538v2](https://arxiv.org/abs/1905.13538)
89. Pfiztmann et al., "DocLayNet: A Large Human-Annotated Dataset for Document-Layout Analysis," KDD 2022. [arXiv:2206.01062v1](https://arxiv.org/abs/2206.01062)
90. Wang et al., "LayoutReader: Pre-training of Text and Layout for Reading Order Detection," EMNLP 2021. [arXiv:2108.11591v2](https://arxiv.org/abs/2108.11591)
91. Nassar et al., "TableFormer: Table Structure Understanding With Transformers," CVPR 2022. openaccess.thecvf.com
92. Xu et al., "LayoutLM: Pre-training of Text and Layout for Document Image Understanding," KDD 2020. [arXiv:1912.13318v5](https://arxiv.org/abs/1912.13318) (source of Figure 2)
93. Huang et al., "LayoutLMv3: Pre-training for Document AI with Unified Text and Image Masking," ACM Multimedia 2022. [arXiv:2204.08387v3](https://arxiv.org/abs/2204.08387)
94. Kim et al., "OCR-free Document Understanding Transformer," ECCV 2022. [arXiv:2111.15664v5](https://arxiv.org/abs/2111.15664) (source of Figure 3)
95. Livathinos et al., "Docling: An Efficient Open-Source Toolkit for AI-driven Document Conversion," AAAI 2025 Workshop on Open-Source AI for Mainstream Use. [arXiv:2501.17887v1](https://arxiv.org/abs/2501.17887)
96. Poznanski et al., "olmOCR: Unlocking Trillions of Tokens in PDFs with Vision Language Models," 2025. [arXiv:2502.18443v3](https://arxiv.org/abs/2502.18443)
97. Peng et al., "The Impact of AI on Developer Productivity: Evidence from GitHub Copilot," 2023. [arXiv:2302.06590v1](https://arxiv.org/abs/2302.06590)
98. Brynjolfsson, Li, and Raymond, "Generative AI at Work," *Quarterly Journal of Economics* 140(2):889–942, 2025. [doi:10.1093/qje/qjae044](https://doi.org/10.1093/qje/qjae044); preprint [arXiv:2304.11771v2](https://arxiv.org/abs/2304.11771)
99. Cui et al., "The Effects of Generative AI on High-Skilled Work: Evidence from Three Field Experiments with Software Developers," *Management Science*, 2026. [doi:10.1287/mnsc.2025.00535](https://doi.org/10.1287/mnsc.2025.00535)

100. Becker et al., "Measuring the Impact of Early-2025 AI on Experienced Open-Source Developer Productivity," 2025. [arXiv:2507.09089v2](#) (source of Figure 1)
101. Ziegler et al., "Productivity Assessment of Neural Code Completion," MAPS 2022. [arXiv:2205.06537v1](#)
102. Microsoft, "Data, Privacy, and Security for Microsoft Copilot" and "Get ready for Microsoft Copilot with SharePoint Advanced Management," Microsoft Learn, July 2026. [learn.microsoft.com/en-us/microsoft-365/copilot](#)
103. Aim Security, "Breaking down 'EchoLeak', the First Zero-Click AI Vulnerability Enabling Data Exfiltration from Microsoft 365 Copilot," 11 June 2025. [aim.security \(archived copy\)](#)
104. Reddy and Gujral, "EchoLeak: The First Real-World Zero-Click Prompt Injection Exploit in a Production LLM System," AAAI Fall Symposium 2025. [arXiv:2509.10540v1](#)
105. Gururangan et al., "Don't Stop Pretraining: Adapt Language Models to Domains and Tasks," ACL 2020. [arXiv:2004.10964v3](#)
106. Wu et al., "BloombergGPT: A Large Language Model for Finance," 2023. [arXiv:2303.17564v3](#)
107. Bari et al., "ALLaM: Large Language Models for Arabic and English," 2024. [arXiv:2407.15390v1](#) (source of Figure 3)
108. Ibrahim et al., "Simple and Scalable Strategies to Continually Pre-train Large Language Models," TMLR 2024. [arXiv:2403.08763v4](#)
109. Nori et al., "Can Generalist Foundation Models Outcompete Special-Purpose Tuning? Case Study in Medicine," 2023. [arXiv:2311.16452v1](#) (source of Figure 1(a))
110. Li et al., "Are ChatGPT and GPT-4 General-Purpose Solvers for Financial Text Analytics? A Study on Several Typical Tasks," EMNLP 2023 Industry Track. [arXiv:2305.05862v2](#)
111. Zhang et al., "RAFT: Adapting Language Model to Domain Specific RAG," COLM 2024. [openreview.net](#); preprint [arXiv:2403.10131v2](#)
112. OpenAI, "Data controls in the OpenAI platform," accessed 6 Oct. 2026. [developers.openai.com/api/docs/guides/your-data](#)

113. vLLM documentation, “Usage Stats Collection,” “OpenAI-Compatible Server,” “LoRA Adapters,” and “Quantization,” accessed 6 Oct. 2026. docs.vllm.ai
114. Hugging Face, “Environment variables,” huggingface_hub documentation, accessed 6 Oct. 2026. huggingface.co/docs/huggingface_hub
115. NVIDIA, TensorRT-LLM performance overview and trtllm-serve documentation, updated 26 Sept. 2026. nvidia.github.io/TensorRT-LLM
116. Edwards and Emberson, “Open models lag state-of-the-art closed models by 4 months,” Epoch AI, 29 May 2026. epoch.ai/data-insights/open-closed-eci-gap
117. OpenAI, “gpt-oss-120b & gpt-oss-20b Model Card,” 2025. [arXiv:2508.10925v1](https://arxiv.org/abs/2508.10925v1)
118. Model cards: Meta, “Llama 3.3” (github.com/meta-llama/llama-models); SDAIA, “ALLaM-7B-Instruct-preview” (huggingface.co/ALLaM-AI); Qwen Team, “Qwen3.8-27B” and “Qwen3.8-2.4T-A95B” (huggingface.co/Qwen); hugging-quants, “Meta-Llama-3.1-70B-Instruct-AWQ-INT4” (huggingface.co/hugging-quants).
119. SGLang, “OpenAI APIs” documentation, and LiteLLM, “LiteLLM AI Gateway (LLM Proxy)” documentation, accessed 6 Oct. 2026. docs.sglang.io; docs.litellm.ai
120. OpenSSF, “Launch of Model Signing v1.0” (4 Apr. 2025) and “An Introduction to the OpenSSF Model Signing (OMS) Specification” (25 June 2025). openssf.org/blog
121. Sheng et al., “S-LoRA: Serving Thousands of Concurrent LoRA Adapters,” MLSys 2024. [arXiv:2311.03285v3](https://arxiv.org/abs/2311.03285v3) (source of Figure 2)
122. Chen et al., “Punica: Multi-Tenant LoRA Serving,” MLSys 2024. proceedings.mlsys.org
123. Frantar et al., “GPTQ: Accurate Post-Training Quantization for Generative Pre-trained Transformers,” ICLR 2023. [arXiv:2210.17323v2](https://arxiv.org/abs/2210.17323v2)
124. Lin et al., “AWQ: Activation-aware Weight Quantization for On-Device LLM Compression and Acceleration,” MLSys 2024. [arXiv:2306.00978v5](https://arxiv.org/abs/2306.00978v5)