

Value-Based Methods: Q-Learning to DQN

Naeemullah Khan

naeemullah.khan@kaust.edu.sa



جامعة الملك عبد الله
للعلوم والتقنية

King Abdullah University of
Science and Technology

KAUST Academy
King Abdullah University of Science and Technology

October 4, 2026

1. Learning Outcomes
2. Recap
3. SARSA
4. Training the Q-network
5. DQN Algorithm
6. Double DQN
7. Case Study: Playing Atari Games
8. Summary
9. References

By the end of this session, you will be able to:

- ▶ Understand and implement **Deep Q-Networks (DQN)**.
- ▶ Grasp the **SARSA algorithm** and how it contrasts with Q-learning.
- ▶ Analyze the **differences between off-policy and on-policy learning**.
- ▶ Understand **Experience Replay** and **Target Networks** as stabilisation techniques.
- ▶ Understand **Double DQN** and the problem of overestimation bias.
- ▶ Describe the key extensions in the **DQN family**: Dueling, Prioritized Replay, Distributional RL, and Rainbow.

Q-Learning Problem	DQN Solution
Curse of dimensionality	Neural network function approximator
Correlated samples $\Rightarrow$ instability	Experience Replay
Moving target $\Rightarrow$ divergence	Target Network
Overestimation bias	Double DQN

Reinforcement Learning: **SARSA**

► **SARSA** stands for:

State
→ Action
→ Reward
→ State'
→ Action'

► **On-policy learning:** Learns Q-values by following the current behavior policy.

► **Update Rule:**

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma Q(s', a') - Q(s, a)]$$

- ▶ Uses the action the agent actually takes, not just the best possible one.
- ▶ Learns from the agent's real experience, not hypothetical alternatives.
- ▶ Tends to be safer—useful in situations where mistakes are costly.
- ▶ Helps the agent improve its actual behavior step by step.

- ▶ The SARSA algorithm is a slight variation of the Q-Learning algorithm.
- ▶ Q-Learning is **off-policy**: the Bellman update uses $\max_{a'} Q(s', a')$, which asks “what is the *best possible* action in s' ?”—regardless of what the agent actually does.
- ▶ SARSA is **on-policy**: the update uses $Q(s', a')$ where a' is the action the current policy actually selects, so it learns from real experience, not hypothetical alternatives.

- ▶ The SARSA algorithm is a slight variation of the Q-Learning algorithm.
- ▶ Q-Learning is **off-policy**: the Bellman update uses $\max_{a'} Q(s', a')$, which asks “what is the *best possible* action in s' ?”—regardless of what the agent actually does.
- ▶ SARSA is **on-policy**: the update uses $Q(s', a')$ where a' is the action the current policy actually selects, so it learns from real experience, not hypothetical alternatives.

$$\text{Q-Learning: } Q(s, a) \leftarrow Q(s, a) + \alpha \left[r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right]$$

- ▶ The SARSA algorithm is a slight variation of the Q-Learning algorithm.
- ▶ Q-Learning is **off-policy**: the Bellman update uses $\max_{a'} Q(s', a')$, which asks “what is the *best possible* action in s' ?”—regardless of what the agent actually does.
- ▶ SARSA is **on-policy**: the update uses $Q(s', a')$ where a' is the action the current policy actually selects, so it learns from real experience, not hypothetical alternatives.

$$\textbf{Q-Learning: } Q(s, a) \leftarrow Q(s, a) + \alpha \left[r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right]$$

$$\textbf{SARSA: } Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma Q(s', a') - Q(s, a)]$$

- ▶ The update equation for SARSA depends on the current state, current action, reward obtained, next state, and next action.
- ▶ This observation led to the naming of the learning technique as **SARSA**, which stands for **State-Action-Reward-State-Action**, symbolizing the tuple (s, a, r, s', a') .

- ▶ The update equation for SARSA depends on the current state, current action, reward obtained, next state, and next action.
- ▶ This observation led to the naming of the learning technique as **SARSA**, which stands for **State-Action-Reward-State-Action**, symbolizing the tuple (s, a, r, s', a') .
- ▶ Similar to DQN, there is also a Deep SARSA variant.

Initialize $Q(s, a)$ arbitrarily for all s, a ; set step size $\alpha \in (0, 1]$ and discount γ

for each episode **do**

 Observe initial state s

 Choose a from s using a policy derived from Q (e.g. ϵ -greedy)

repeat

 Take action a ; observe reward r and next state s'

 Choose a' from s' using a policy derived from Q (e.g. ϵ -greedy)

$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma Q(s', a') - Q(s, a)]$

$s \leftarrow s'$; $a \leftarrow a'$

until s is terminal

end

Algorithm: SARSA (on-policy TD control)

⁰Sutton & Barto, *Reinforcement Learning: An Introduction* (2018), Ch. 6. The next action a' is chosen *before* the update—this is what makes SARSA on-policy.

Q-Learning vs SARSA: Comparison Table

Feature	Q-Learning	SARSA
Policy Type	Off-policy	On-policy
Learns From	Others' or past experience	Own current experience
Target	$\max_{a'} Q(s', a')$	$Q(s', a')$ actually taken
Exploration Aware	No	Yes
Risk Sensitivity	More aggressive	More conservative
Use Case	Goal-seeking behavior	Risk-aware behavior

The on-policy vs. off-policy distinction is one of the most important axes in RL algorithm design and will appear in every remaining topic of this course.

Reinforcement Learning: **Training the Q-network**

- ▶ Learning from batches of consecutive samples is problematic:
 - Samples are correlated $\Rightarrow$ inefficient learning.
 - The current Q-network parameters determine the next training samples (e.g., if the maximizing action is to move left, training samples will be dominated by transitions from the left-hand side), which can lead to bad feedback loops.

- ▶ Learning from batches of consecutive samples is problematic:
 - Samples are correlated $\Rightarrow$ inefficient learning.
 - The current Q-network parameters determine the next training samples (e.g., if the maximizing action is to move left, training samples will be dominated by transitions from the left-hand side), which can lead to bad feedback loops.
- ▶ These problems can be addressed using experience replay:
 - Maintain a replay memory buffer of transitions (s, a, r, s') as episodes are played.
 - Train the Q-network on random minibatches of transitions sampled from the replay memory, instead of using consecutive samples.
- ▶ Note: the replay buffer is a general off-policy tool—it is not specific to DQN.

```
Initialize replay memory  $\mathcal{D}$  to capacity  $N$ 
Initialize action-value network  $Q$  with random weights  $\theta$ 
for  $episode = 1$  to  $M$  do
    Observe initial state  $s$ 
    for  $t = 1$  to  $T$  do
        With probability  $\varepsilon$  select a random action  $a$ ; otherwise  $a = \arg \max_{a'} Q(s, a'; \theta)$ 
        Execute  $a$ ; observe reward  $r$  and next state  $s'$ 
        Store transition  $(s, a, r, s')$  in  $\mathcal{D}$ ;  $s \leftarrow s'$ 
        Sample a random minibatch of transitions  $(s_i, a_i, r_i, s'_i)$  from  $\mathcal{D}$ 
        Set  $y_i = \begin{cases} r_i & \text{if } s'_i \text{ is terminal} \\ r_i + \gamma \max_{a'} Q(s'_i, a'; \theta) & \text{otherwise} \end{cases}$ 
        Take a gradient descent step on  $(y_i - Q(s_i, a_i; \theta))^2$  w.r.t.  $\theta$ 
    end
end
```

Algorithm: Deep Q-Network (DQN) with Experience Replay

⁰ ε -greedy exploration—covered with tabular Q-learning. ε is annealed from high to low over training.
Prediction and target share the *same* weights θ —the instability this causes motivates the next slide. ▶ ◀ ≡ ≡ ≡

- ▶ DQN, as stated above, will have learning problems because the target and the prediction are not independent as they both rely on the same network.
- ▶ It's like a dog chasing its own tail.

- ▶ DQN, as stated above, will have learning problems because the target and the prediction are not independent as they both rely on the same network.
- ▶ It's like a dog chasing its own tail.
- ▶ **Solution:** Use two separate Q -value estimators, each of which is used to update the other.
- ▶ The target values are calculated using a target Q -network. The target Q -network's parameters are updated to the current networks every C time steps.
- ▶ Target network prevents the network from spiraling around.

```
Initialize replay memory  $\mathcal{D}$  to capacity  $N$ 
Initialize online network  $Q$  with random weights  $\theta$ 
Initialize target network  $\hat{Q}$  with weights  $\theta^- \leftarrow \theta$ 
for  $episode = 1$  to  $M$  do
    Observe initial state  $s$ 
    for  $t = 1$  to  $T$  do
        With probability  $\epsilon$  select a random action  $a$ ; otherwise  $a = \arg \max_{a'} Q(s, a'; \theta)$ 
        Execute  $a$ ; observe reward  $r$  and next state  $s'$ 
        Store transition  $(s, a, r, s')$  in  $\mathcal{D}$ ;  $s \leftarrow s'$ 
        Sample a random minibatch of transitions  $(s_i, a_i, r_i, s'_i)$  from  $\mathcal{D}$ 

        Set  $y_i = \begin{cases} r_i & \text{if } s'_i \text{ is terminal} \\ r_i + \gamma \max_{a'} \hat{Q}(s'_i, a'; \theta^-) & \text{otherwise} \end{cases}$ 

        Take a gradient descent step on  $(y_i - Q(s_i, a_i; \theta))^2$  w.r.t.  $\theta$ 
        Every  $C$  steps set  $\theta^- \leftarrow \theta$ 
    end
end
```

Algorithm: DQN with a Target Network

⁰Mnih et al. (2015), **Human-level control through deep reinforcement learning**, *Nature*. The target y_i now uses the frozen network $\hat{Q}(\cdot; \theta^-)$.

Reinforcement Learning: **Double DQN**

The problem. Even with a target network, the target computation still has a flaw:

$$y = r + \gamma \max_{a'} \hat{Q}(s', a'; \theta^-)$$

The target network both *selects* the best next action (via max) and *scores* it (via $\hat{Q}$) in one step. Taking a max over noisy estimates will almost always pick an overestimated action—and then that overestimate becomes the training target.

The problem. Even with a target network, the target computation still has a flaw:

$$y = r + \gamma \max_{a'} \hat{Q}(s', a'; \theta^-)$$

The target network both *selects* the best next action (via max) and *scores* it (via $\hat{Q}$) in one step. Taking a max over noisy estimates will almost always pick an overestimated action—and then that overestimate becomes the training target.

The fix. Double DQN splits the job between the two networks we already have:

$$y = r + \gamma \hat{Q}\left(s', \arg \max_{a'} Q(s', a'; \theta); \theta^-\right)$$

The **online network** θ picks which action looks best. The **target network** θ^- scores that action. Because their weights differ, their errors do not overlap—overestimation bias is reduced. This is a one-line change to the DQN code.

Both use two networks, which makes them easy to confuse. They solve completely different problems:

	Problem solved	How
Target Network	Unstable, moving target	Freeze weights for C steps
Double DQN	Overestimation bias	Split action selection from evaluation

The two-network structure was already there from the Target Network. Double DQN just changes *how* those two networks are used.

```
Initialize replay memory  $\mathcal{D}$  to capacity  $N$ 
Initialize online network  $Q$  with random weights  $\theta$ 
Initialize target network  $\hat{Q}$  with weights  $\theta^- \leftarrow \theta$ 
for  $episode = 1$  to  $M$  do
    Observe initial state  $s$ 
    for  $t = 1$  to  $T$  do
        With probability  $\epsilon$  select a random action  $a$ ; otherwise  $a = \arg \max_{a'} Q(s, a'; \theta)$ 
        Execute  $a$ ; observe reward  $r$  and next state  $s'$ 
        Store transition  $(s, a, r, s')$  in  $\mathcal{D}$ ;  $s \leftarrow s'$ 
        Sample a random minibatch of transitions  $(s_i, a_i, r_i, s'_i)$  from  $\mathcal{D}$ 
         $a_i^* \leftarrow \arg \max_{a'} Q(s'_i, a'; \theta)$  // select with online net  $\theta$ 
        Set  $y_i = \begin{cases} r_i & \text{if } s'_i \text{ is terminal} \\ r_i + \gamma \hat{Q}(s'_i, a_i^*; \theta^-) & \text{otherwise} \end{cases}$  // evaluate with target net  $\theta^-$ 
        Take a gradient descent step on  $(y_i - Q(s_i, a_i; \theta))^2$  w.r.t.  $\theta$ 
        Every  $C$  steps set  $\theta^- \leftarrow \theta$ 
    end
end
```

Algorithm: Double DQN (DDQN)

⁰Van Hasselt, Guez & Silver (2016), *Deep Reinforcement Learning with Double Q-learning*, AAAI. Only the target line changes vs. DQN: action selection (θ) is split from action evaluation (θ^-).


Case Study: **Playing Atari Games**



- **Objective:** Complete the game with the highest score
- **State:** Raw pixel inputs of the game state
- **Action:** Game controls e.g. Left, Right, Up, Down
- **Reward:** Score increase/decrease at each time step

⁰[Mnih et al. NIPS Workshop 2013; Nature 2015]

$Q(s, a; \theta)$:
neural network
with weights θ




← Last FC layer has 4-d output (if 4 actions), corresponding to $Q(s_t, a_1)$, $Q(s_t, a_2)$, $Q(s_t, a_3)$, $Q(s_t, a_4)$

Number of actions between 4-18 depending on Atari game

Current state s_t : 84x84x4 stack of last 4 frames
(after RGB->grayscale conversion, downsampling, and cropping)

⁰[Mnih et al. NIPS Workshop 2013; Nature 2015]

$Q(s, a; \theta)$:
neural network
with weights θ



← Last FC layer has 4-d
output (if 4 actions),
corresponding to $Q(s_t, a_1)$, $Q(s_t, a_2)$, $Q(s_t, a_3)$,
 $Q(s_t, a_4)$

Number of actions between 4-18
depending on Atari game

Current state s_t : 84x84x4 stack of last 4 frames
(after RGB->grayscale conversion, downsampling, and cropping)

<https://www.youtube.com/watch?v=V1eYniJ0Rnk>

⁰[Mnih et al. NIPS Workshop 2013; Nature 2015]

DQN Family & SARSA: **Summary**

Scenario	Recommended Algorithm
Environment is stochastic	SARSA (conservative)
Maximizing reward is the priority	Q-Learning or DQN
Large state space (e.g., images)	DQN
Limited computational resources	SARSA (simpler model)

DQN works well but has known weaknesses. Each extension below addresses one specific problem. Together they form the DQN family.

- ▶ **Dueling DQN:** Splits Q into $V(s)$ and $A(s, a)$ —learns which states are valuable without evaluating every action.
- ▶ **Prioritized Replay:** Samples transitions proportional to TD error instead of uniformly—prioritises what the network got wrong.
- ▶ **Distributional RL:** Learns the full return distribution $Z(s, a)$, not just the expected value—richer signal, faster learning.
- ▶ **Noisy Nets:** Replaces ε -greedy with learned noise in the weights—the network decides how much to explore.
- ▶ **Rainbow:** Combines all of the above. Represents the practical ceiling of value-based methods.

- ▶ SARSA is an on-policy alternative to Q-learning, using the action the agent actually takes rather than $\max_{a'} Q(s', a')$.
- ▶ The on-policy vs. off-policy distinction is one of the most important axes in RL algorithm design.
- ▶ Deep Q-Learning uses a neural network as a function approximator for Q-values, enabling RL on high-dimensional inputs.
- ▶ Experience replay breaks correlation between consecutive samples; the target network stabilizes the moving target.
- ▶ Double DQN decouples action selection from evaluation to reduce overestimation bias.

- The DQN family extensions (Dueling, Prioritized Replay, Distributional, Noisy Nets, Rainbow) each address a specific weakness of vanilla DQN.

DQN Family & SARSA: **References**

- [1] Mnih, V., Kavukcuoglu, K., Silver, D., et al. (2015).
Human-level control through deep reinforcement learning.
Nature.
- [2] Van Hasselt, H., Guez, A., Silver, D. (2016).
Deep Reinforcement Learning with Double Q-learning.
AAAI.
- [3] Wang, Z., Schaul, T., Hessel, M., et al. (2016).
Dueling Network Architectures for Deep RL.
ICML.

- [4] Schaul, T., Quan, J., Antonoglou, I., Silver, D. (2016).
Prioritized Experience Replay.
ICLR.
- [5] Hessel, M., Modayil, J., Van Hasselt, H., et al. (2018).
Rainbow: Combining Improvements in Deep Reinforcement Learning.
AAAI.
- [6] Bellemare, M. G., Dabney, W., Munos, R. (2017).
A Distributional Perspective on Reinforcement Learning.
ICML.

- [7] Rummery, G. A., & Niranjan, M. (1994).
On-line Q-learning using connectionist systems.
Technical Report.
- [8] Sutton, R. S., & Barto, A. G. (2018).
Reinforcement Learning: An Introduction.
- [9] Berkeley CS285 Deep RL Lectures.
<https://rail.eecs.berkeley.edu/deeprlcourse/>
- [10] David Silver's RL Series.
<http://www0.cs.ucl.ac.uk/staff/d.silver/web/Teaching.html>
- [11] Emma Brunskill, Stanford CS234: Reinforcement Learning.
<https://web.stanford.edu/class/cs234/>

- [12] Fei-Fei Li, Yunzhu Li & Ruohan Gao, Stanford CS231n: Deep Learning for Computer Vision.
<http://cs231n.stanford.edu/>

- [13] Ashwin Rao, Stanford CME241: Foundations of Reinforcement Learning with Applications in Finance.
<https://web.stanford.edu/class/cme241/>

- [14] Jimmy Ba & Bo Wang, UofT CSC413/2516: Neural Networks and Deep Learning.
<https://uoft-csc413.github.io/2022/>